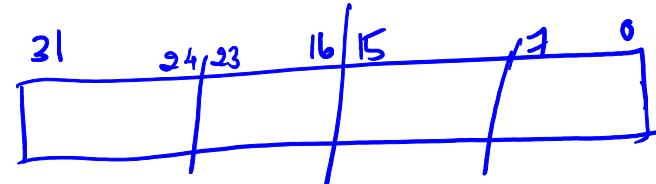


M4 – Parallelism

Outline

- Parallelism
- Flynn's classification
- Vector Processing
- SIMD: Subword Parallelism
- Symmetric Multiprocessors, Distributed Memory Machines
 - Shared Memory Multiprocessing, Message Passing
- Synchronization
 - ISA
- Cache coherence



Flynn's Classification

- **Single instruction stream, single data stream (SISD)**
 - Uniprocessor.

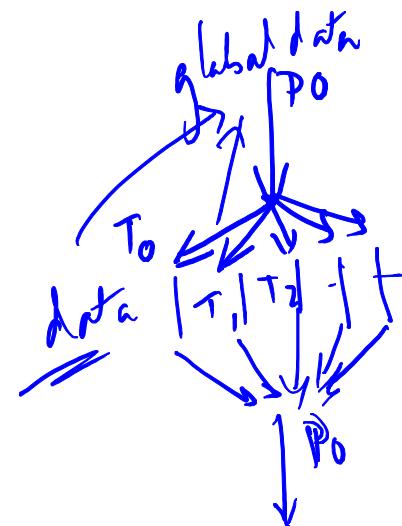
Flynn's Classification

- **Single instruction stream, single data stream (SISD)**
 - Uniprocessor.
- **Single instruction stream, multiple data streams (SIMD)**
 - Data-level parallelism, Subword Parallelism
 - Applying same operations to multiple items of data in parallel
 - Eg. Multimedia extensions, Vector architectures
 - Gaming, 3D, real-time VR, rendering, ...

Flynn's Classification

- Multiple instruction streams, single data stream (**MISD**)
- Multiple instruction streams, multiple data streams (**MIMD**)
 - Thread-level parallelism

Task - level

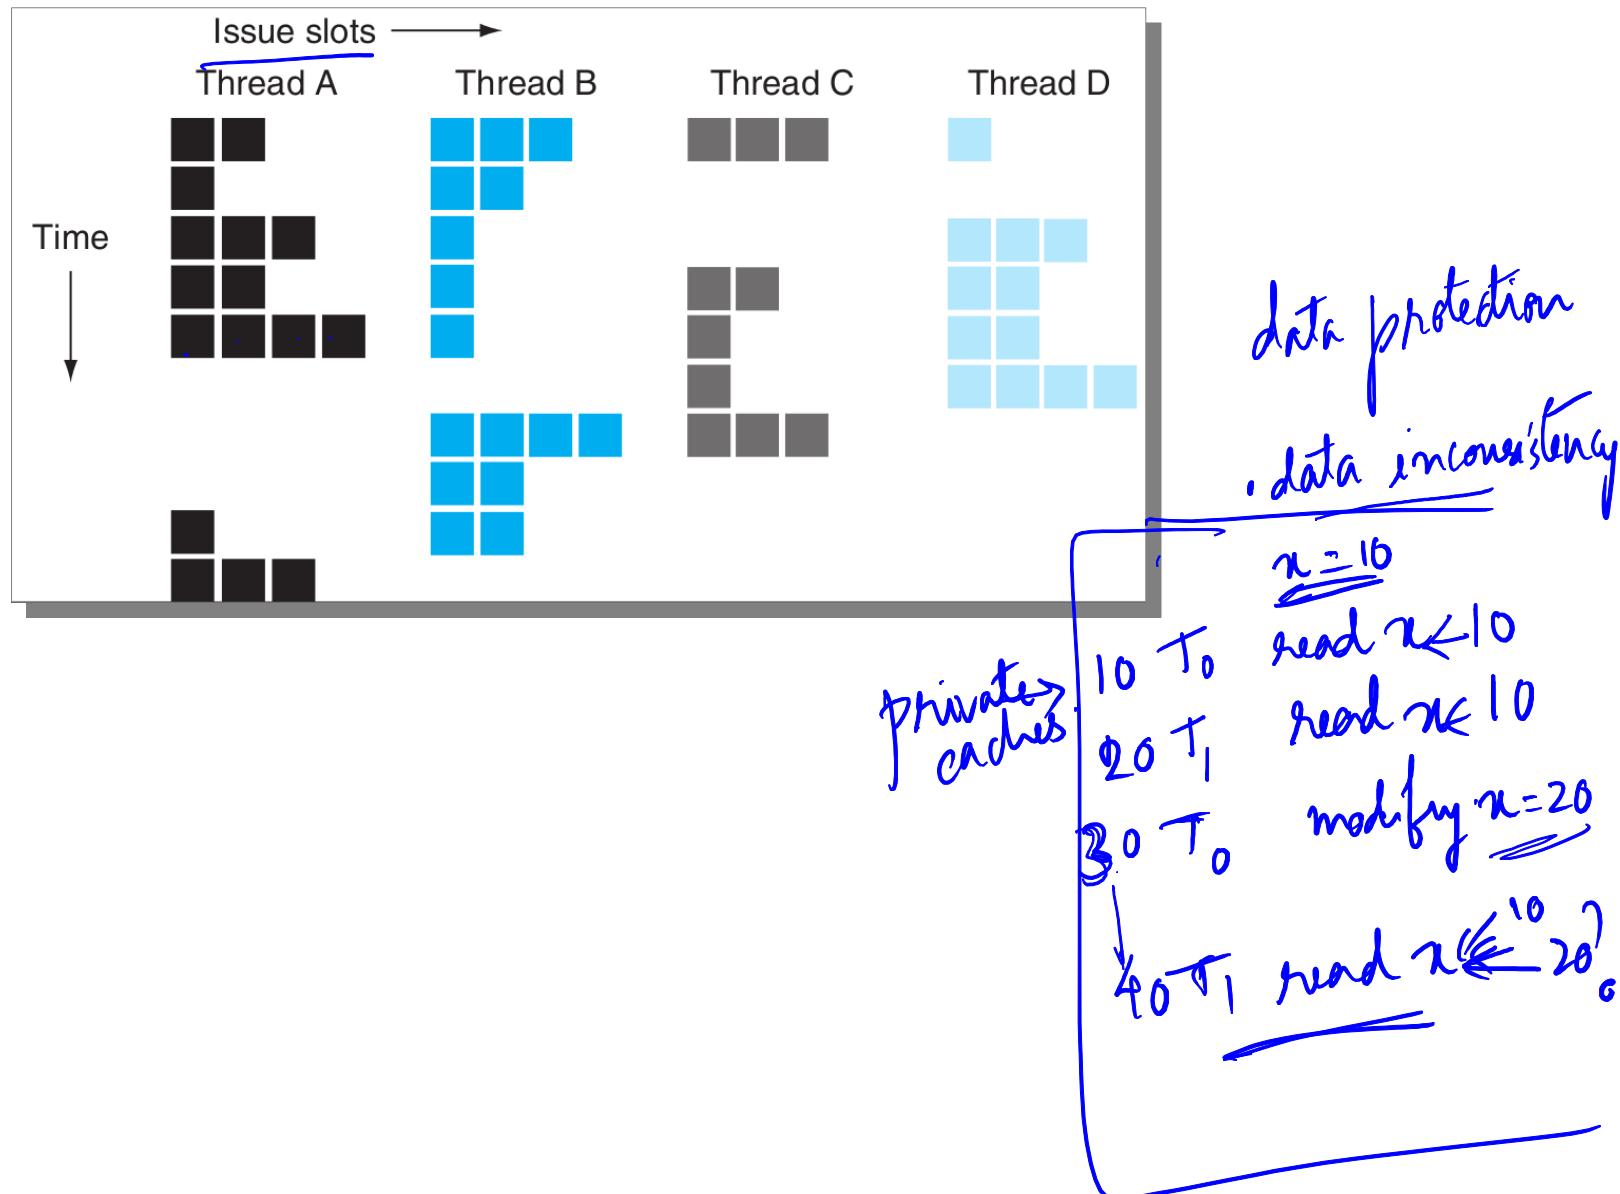


{
• ~~light weight process~~
• has its own data (stack)
• shares the code

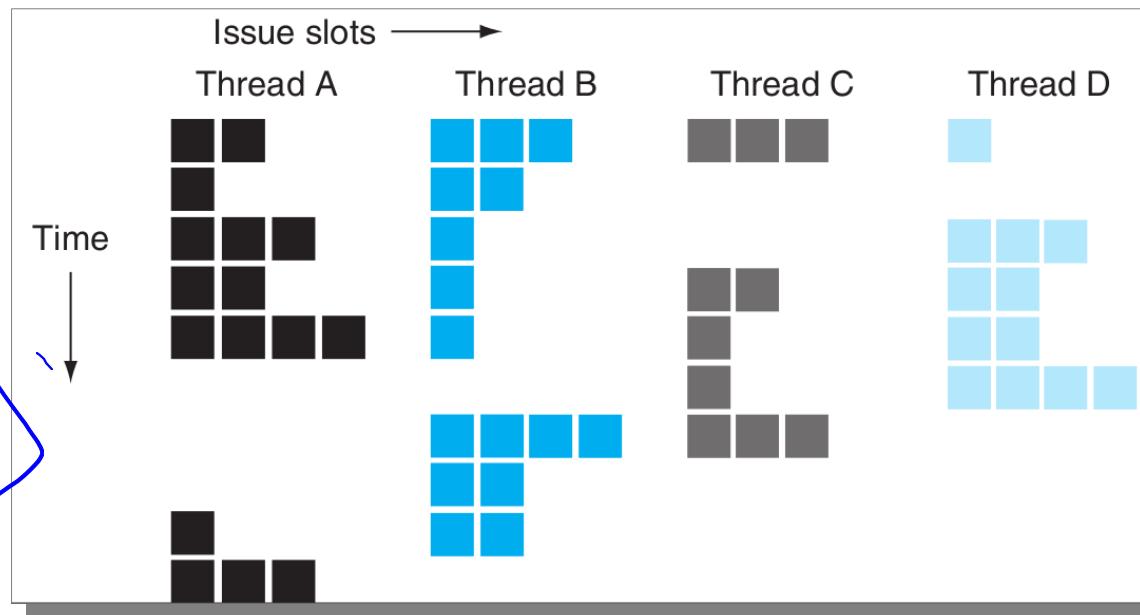
MIMD

Multithreading

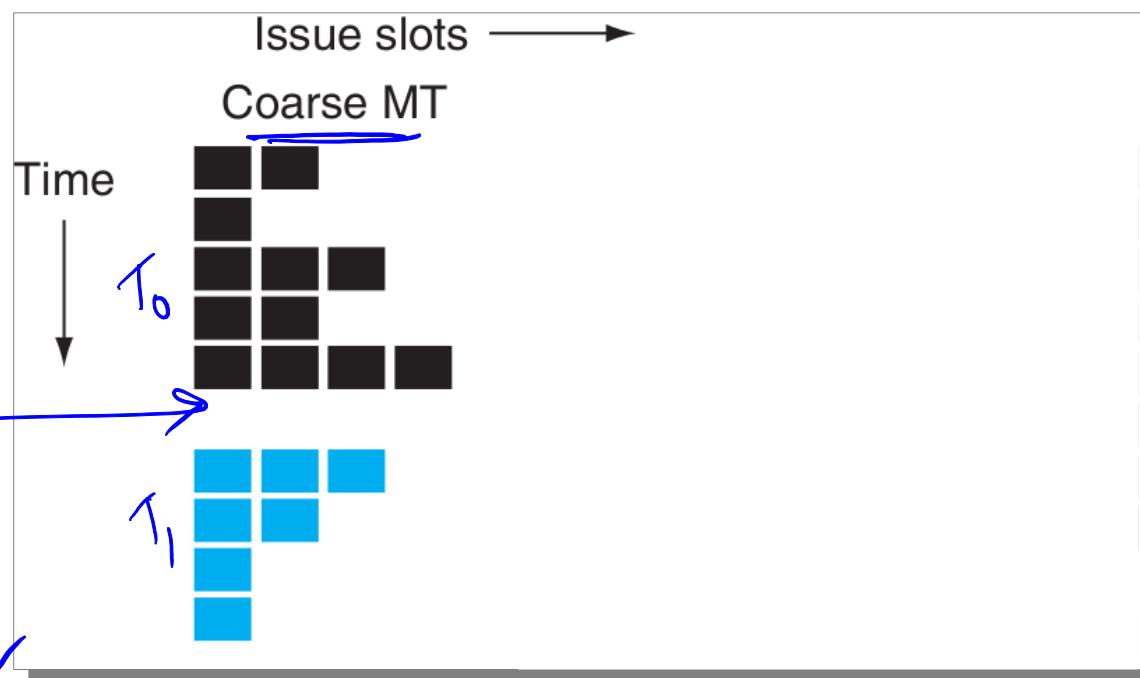
Spawn



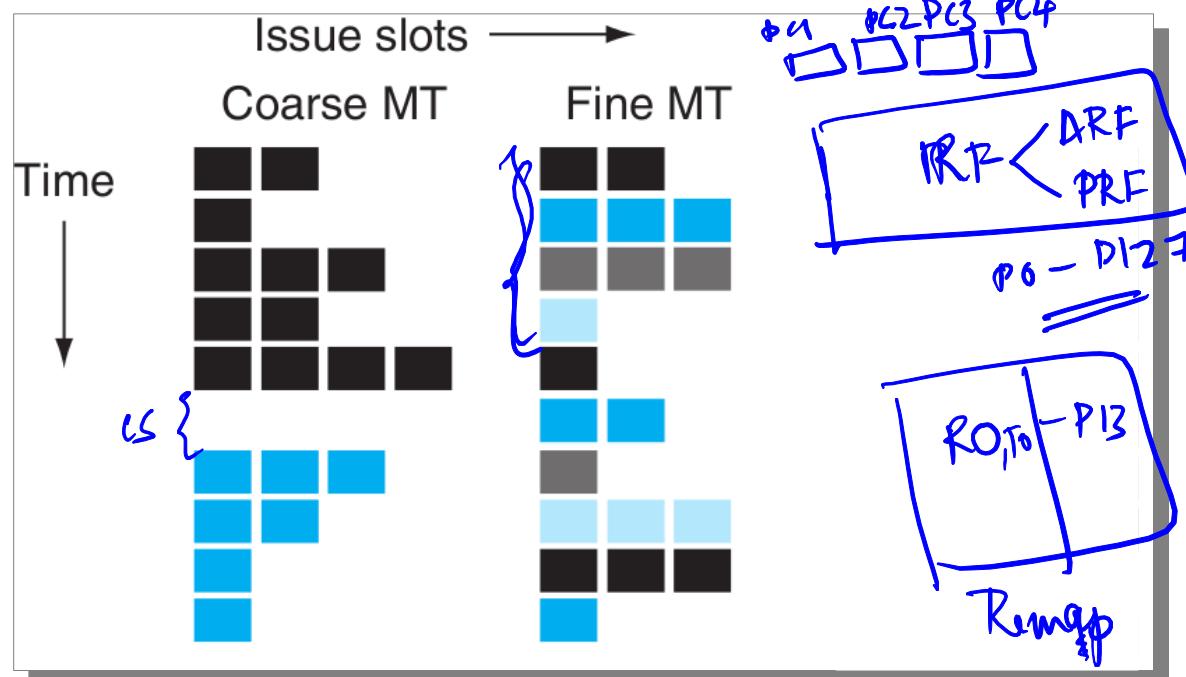
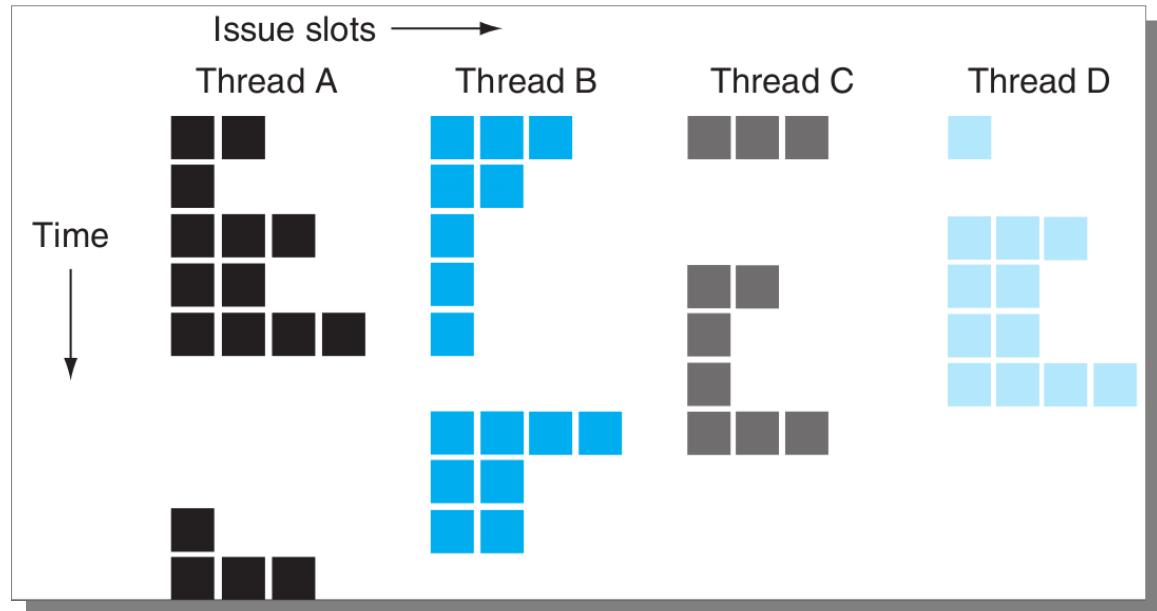
Multithreading



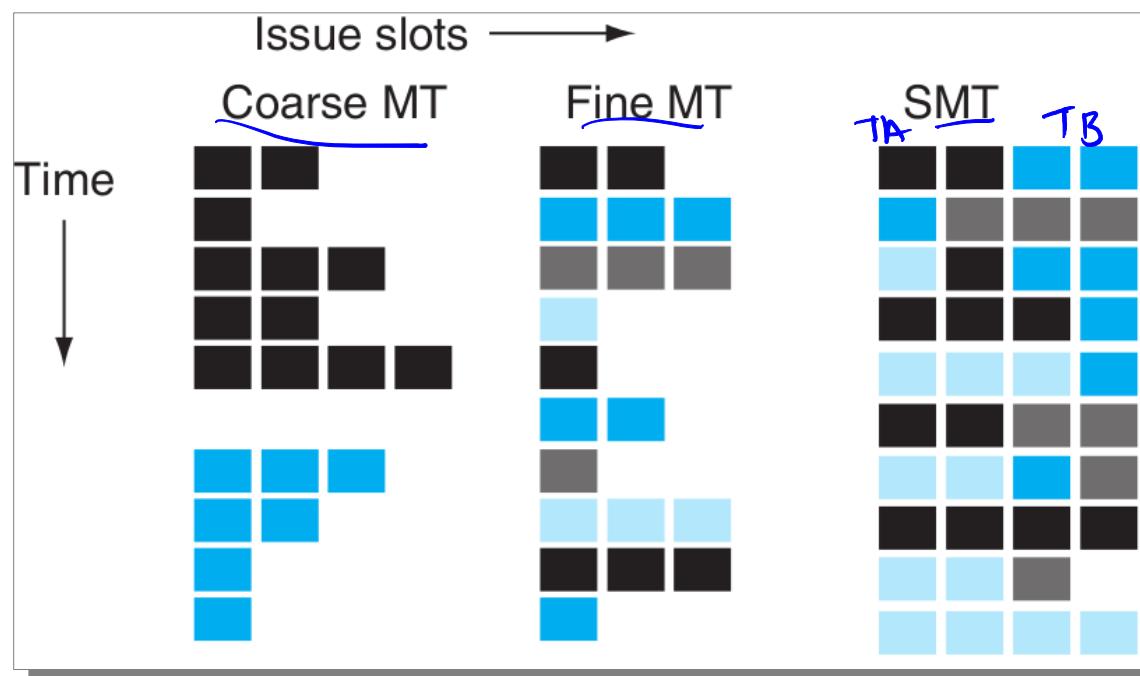
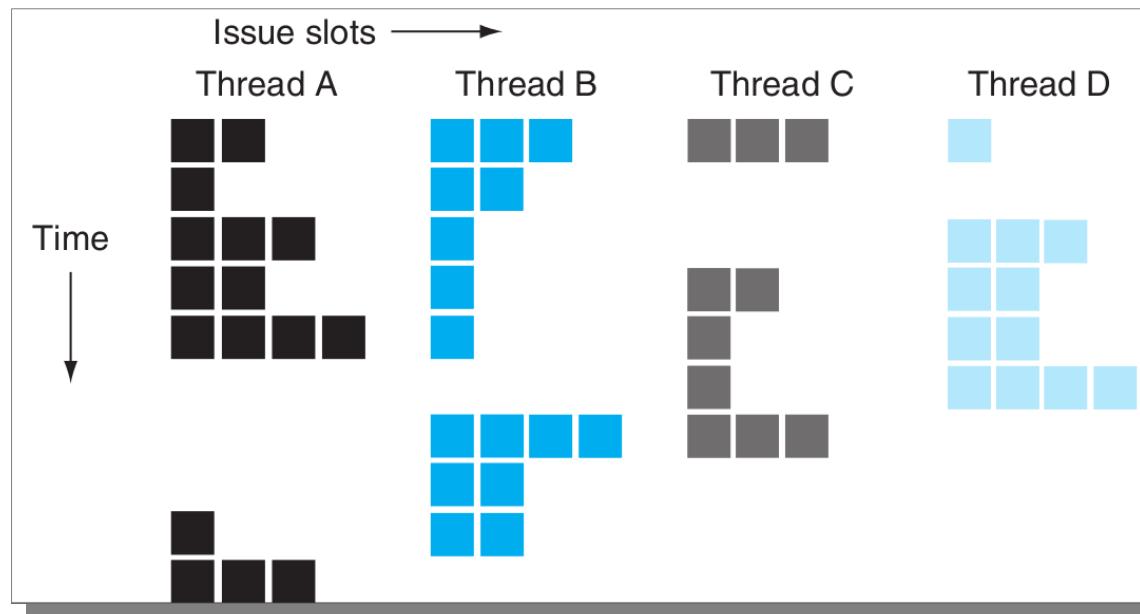
{ update dirty blocks
L2
, equal to dirty contents
0



Multithreading



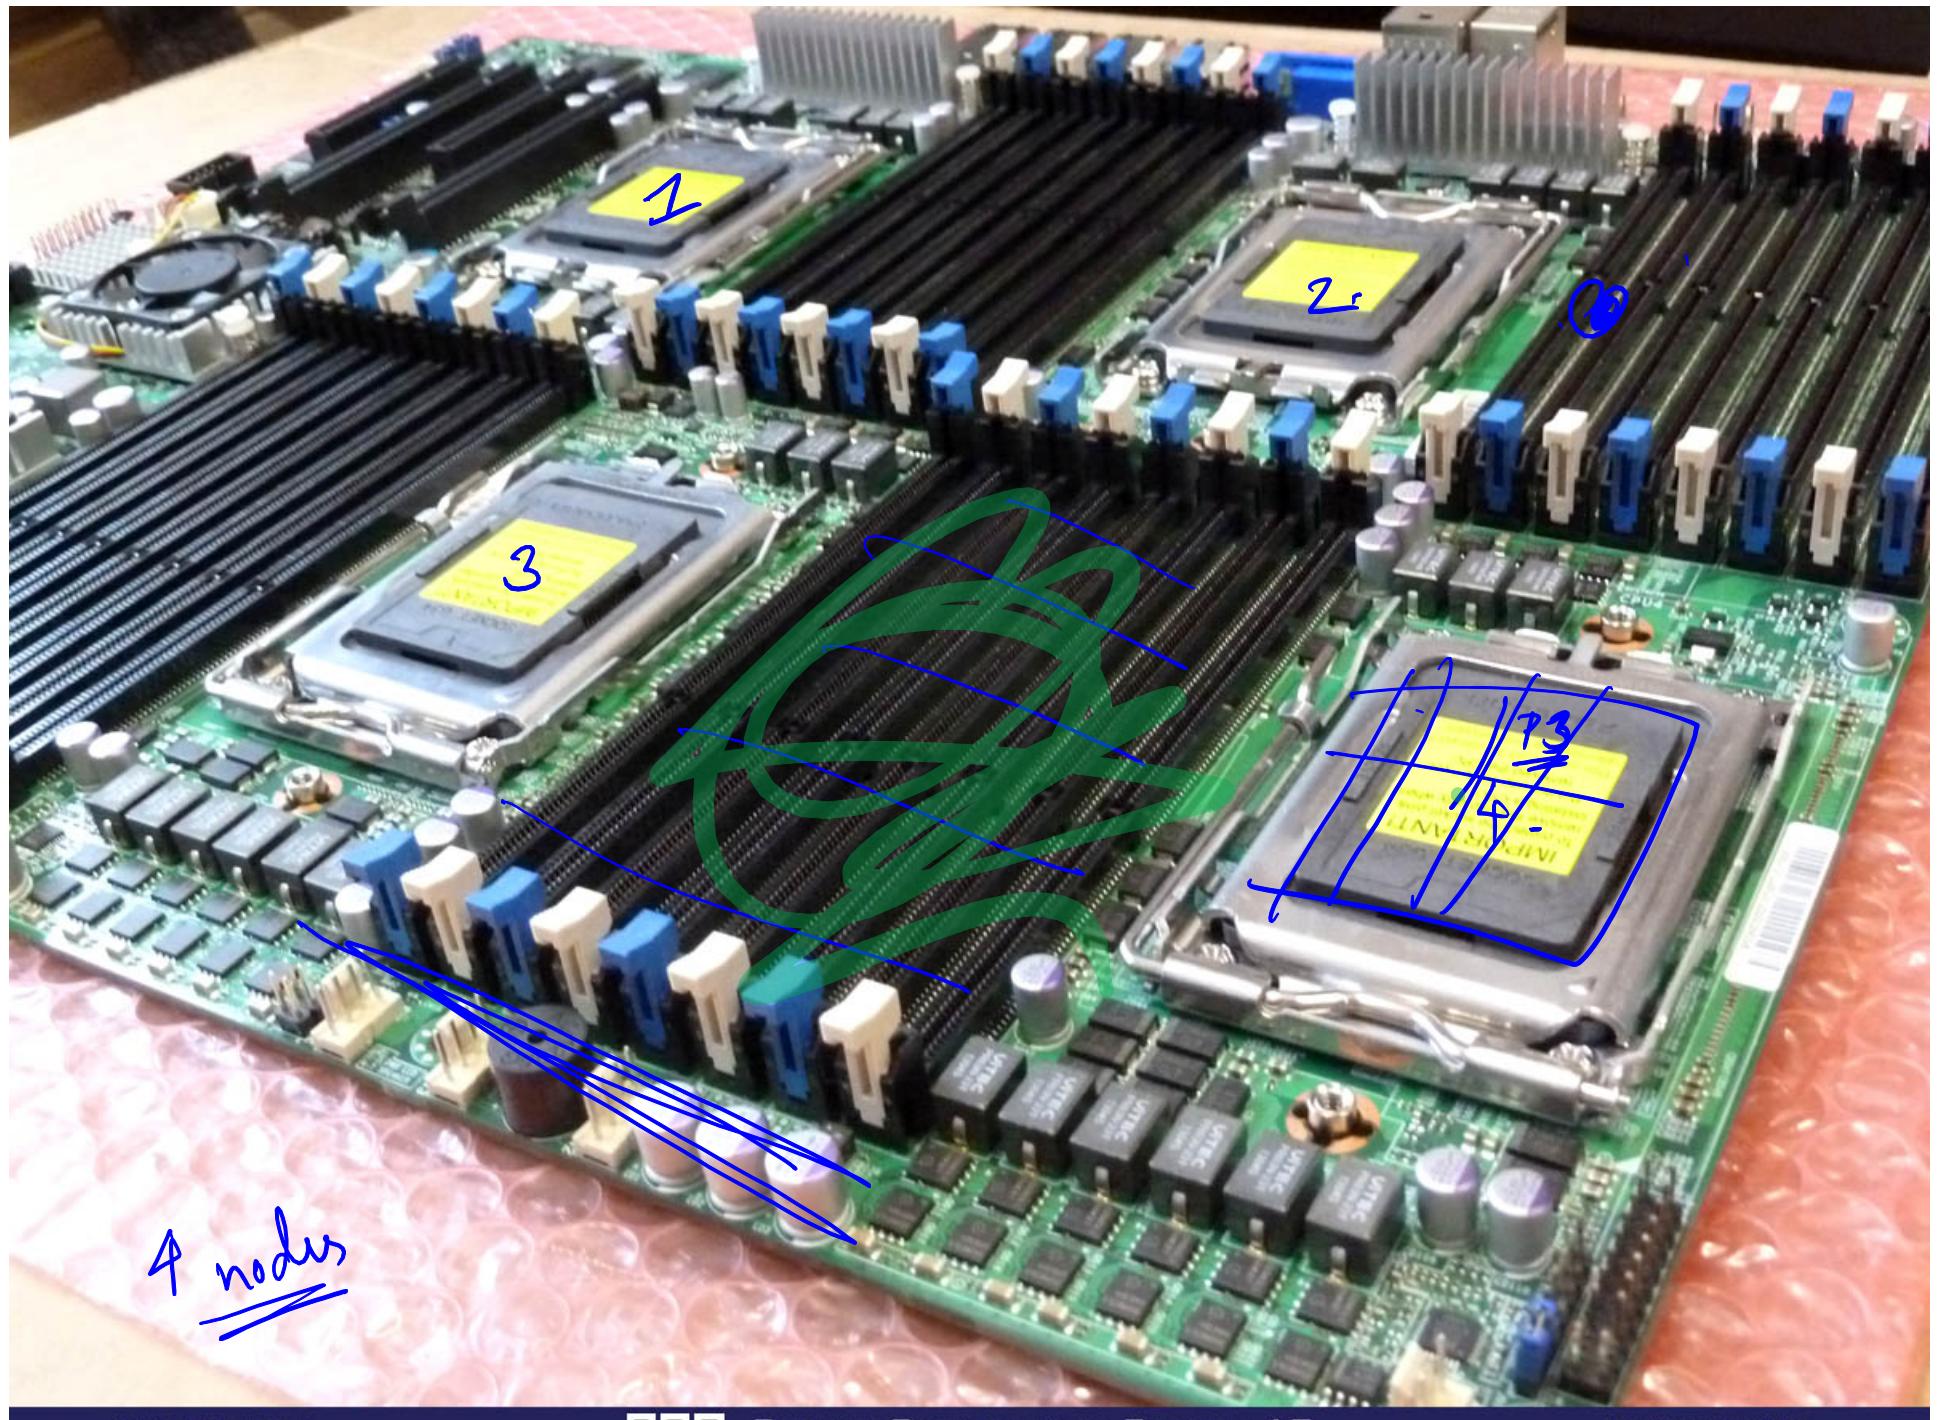
Multithreading



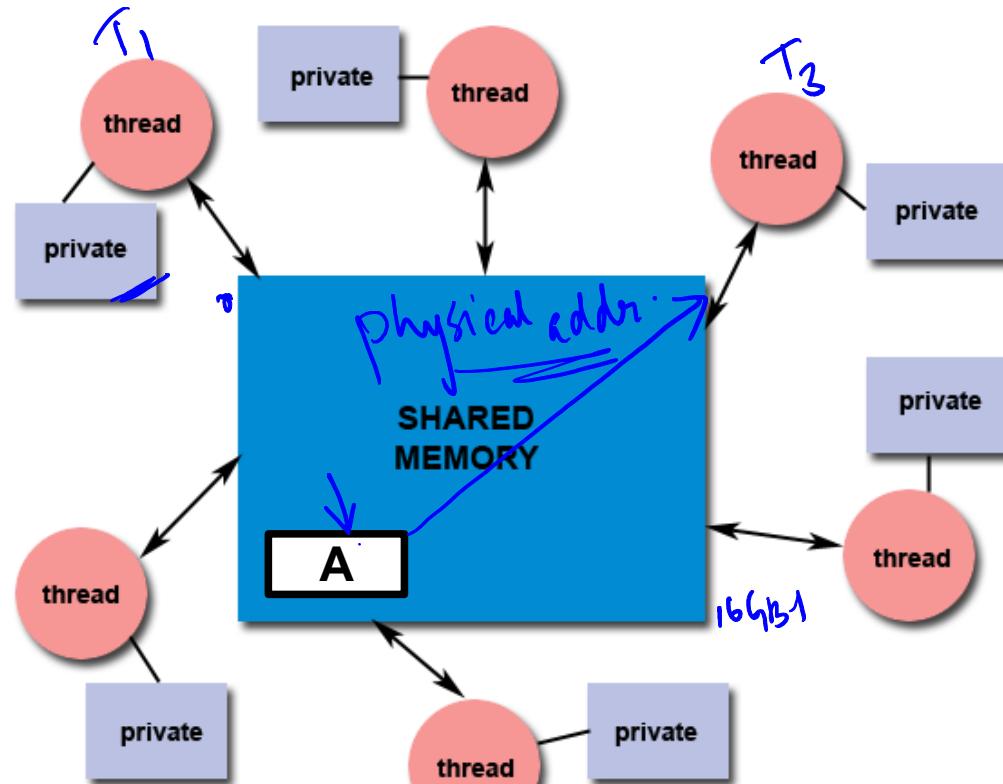
8 core, 2 threads with
Hyperthreading
cycle

Shared vs. Distributed Memory

- Shared – Processors share physical memory space → 1 MB, 1 socket.
- Distributed – Each node has its own physical address space

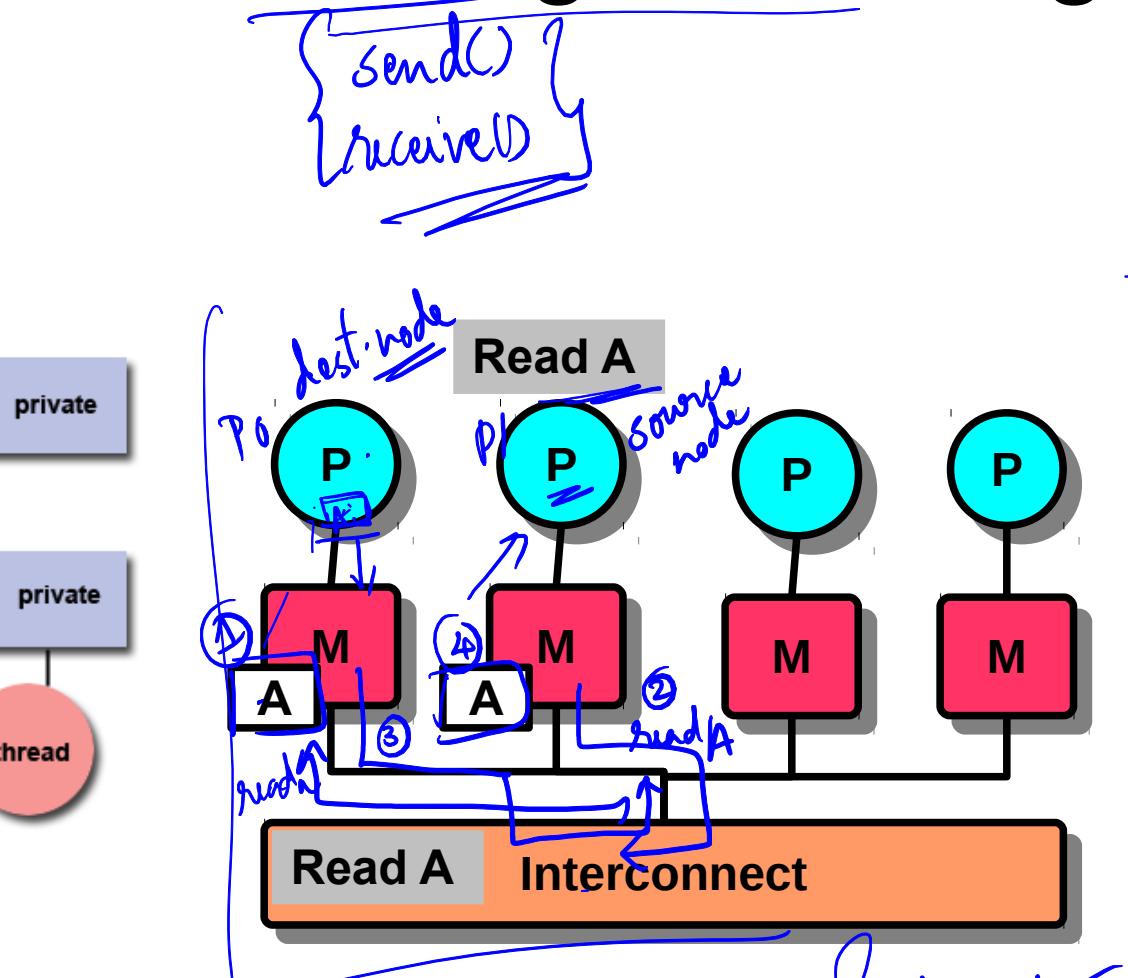


Shared Memory vs. Message Passing



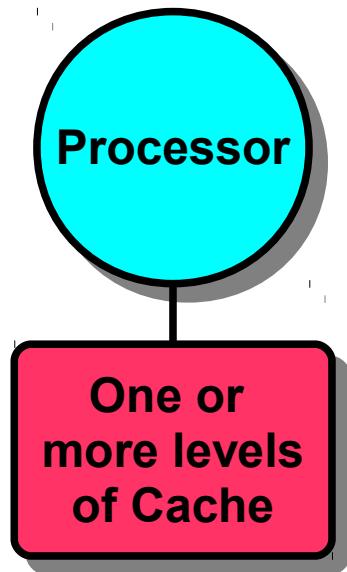
Distributed Mem:

Message Passing

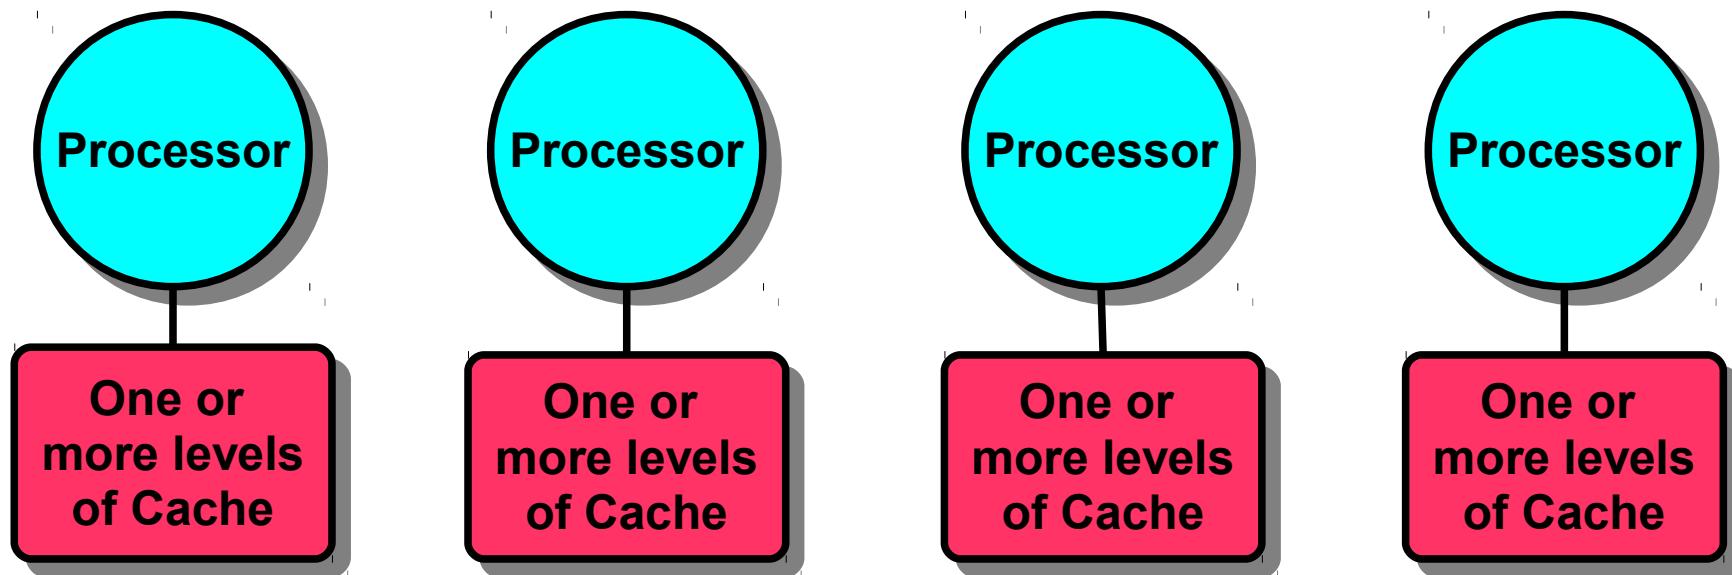


MC: P0: A: write } directory
 P1: read A
 MC: P1 }
 P0: write
 read }

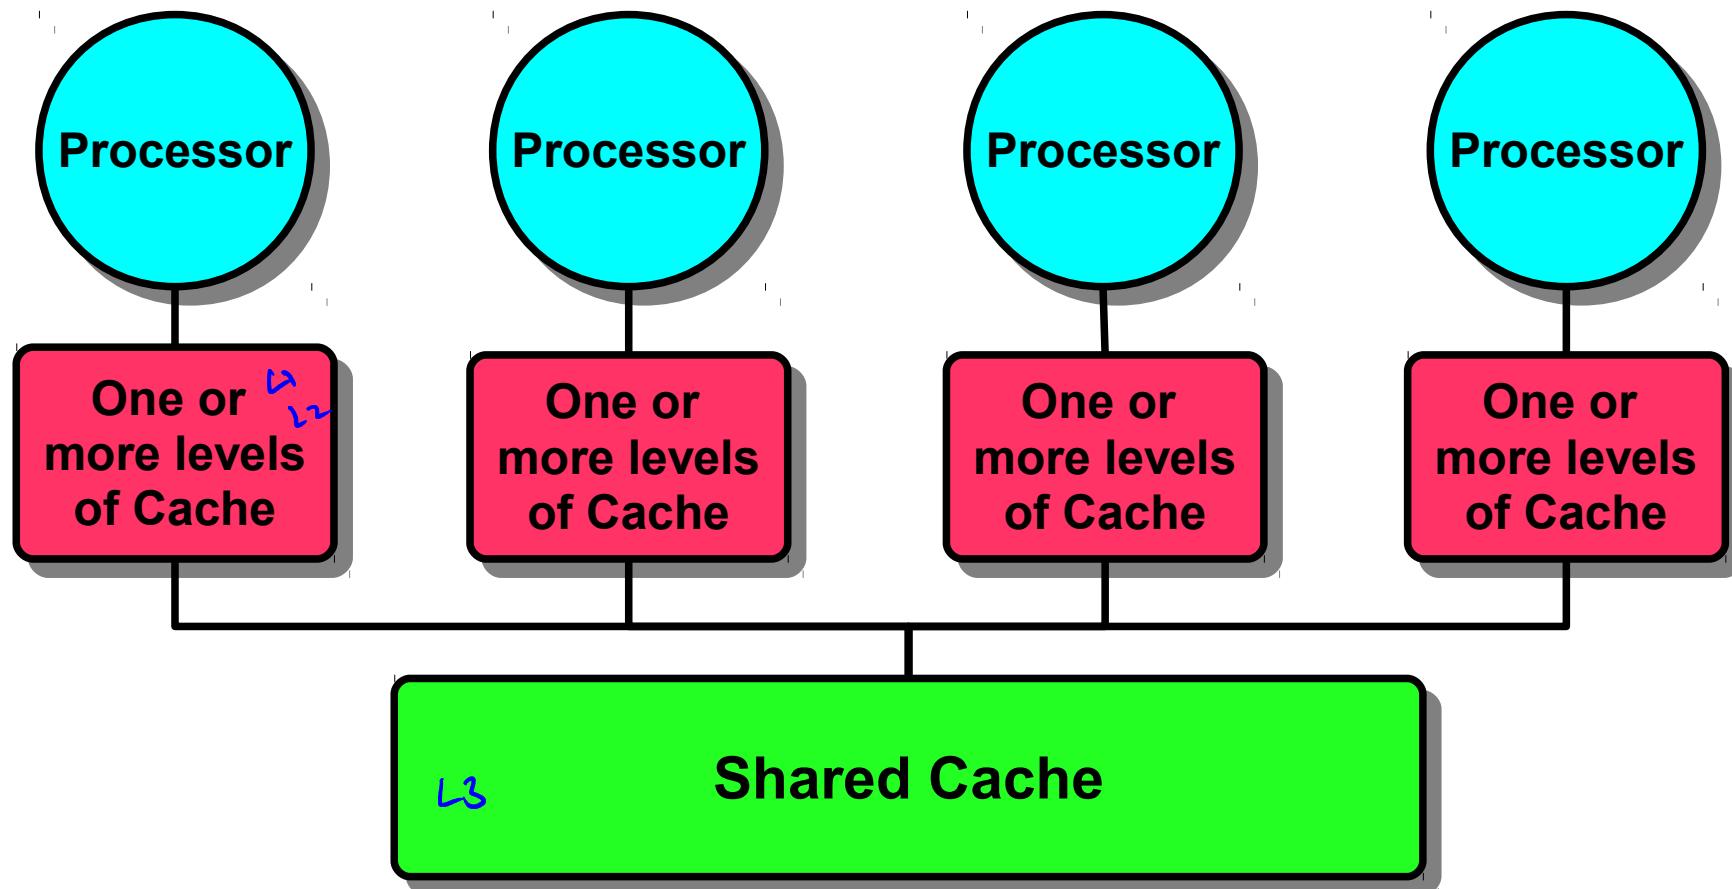
Symmetric Multiprocessor (SMP)



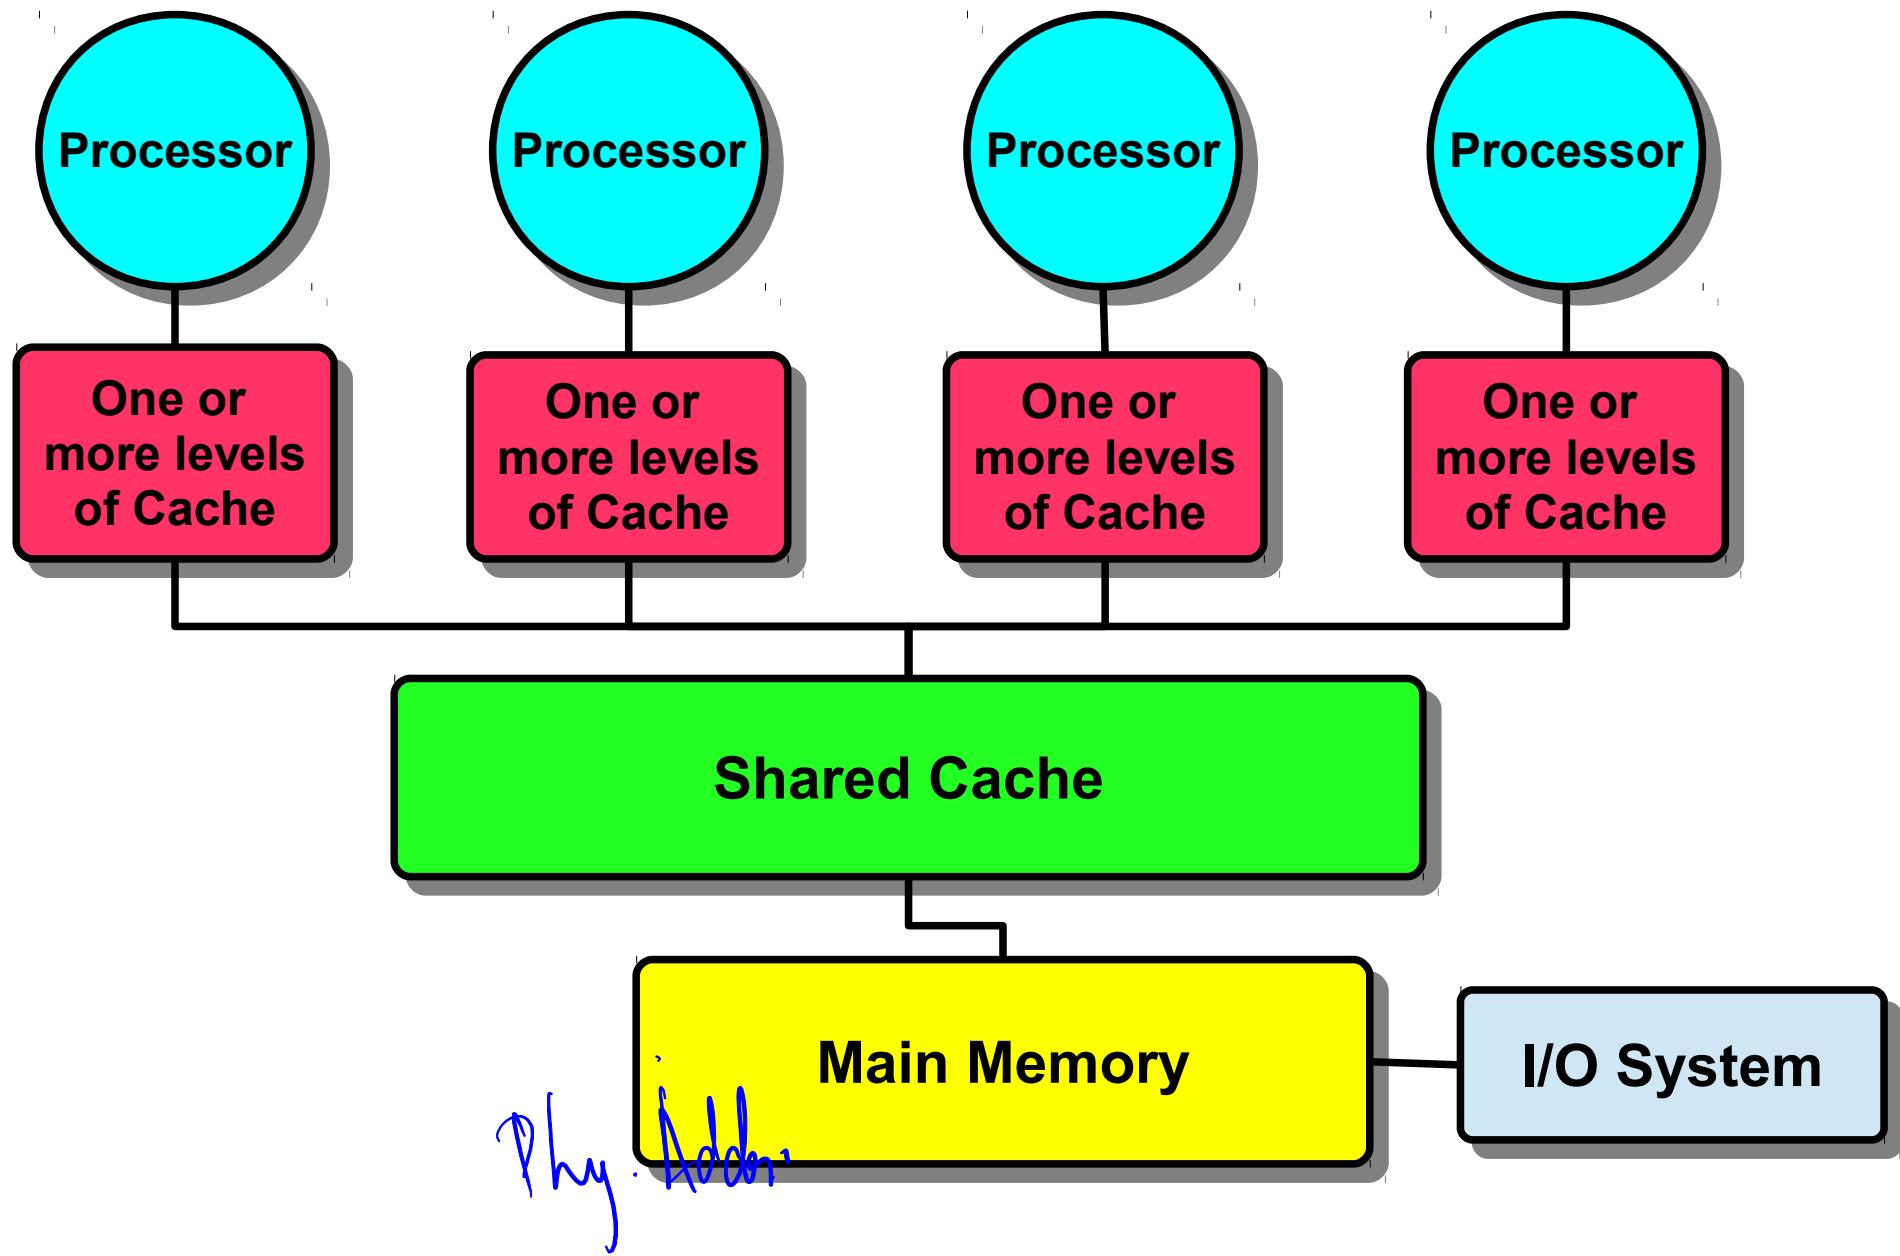
Symmetric Multiprocessor (SMP)



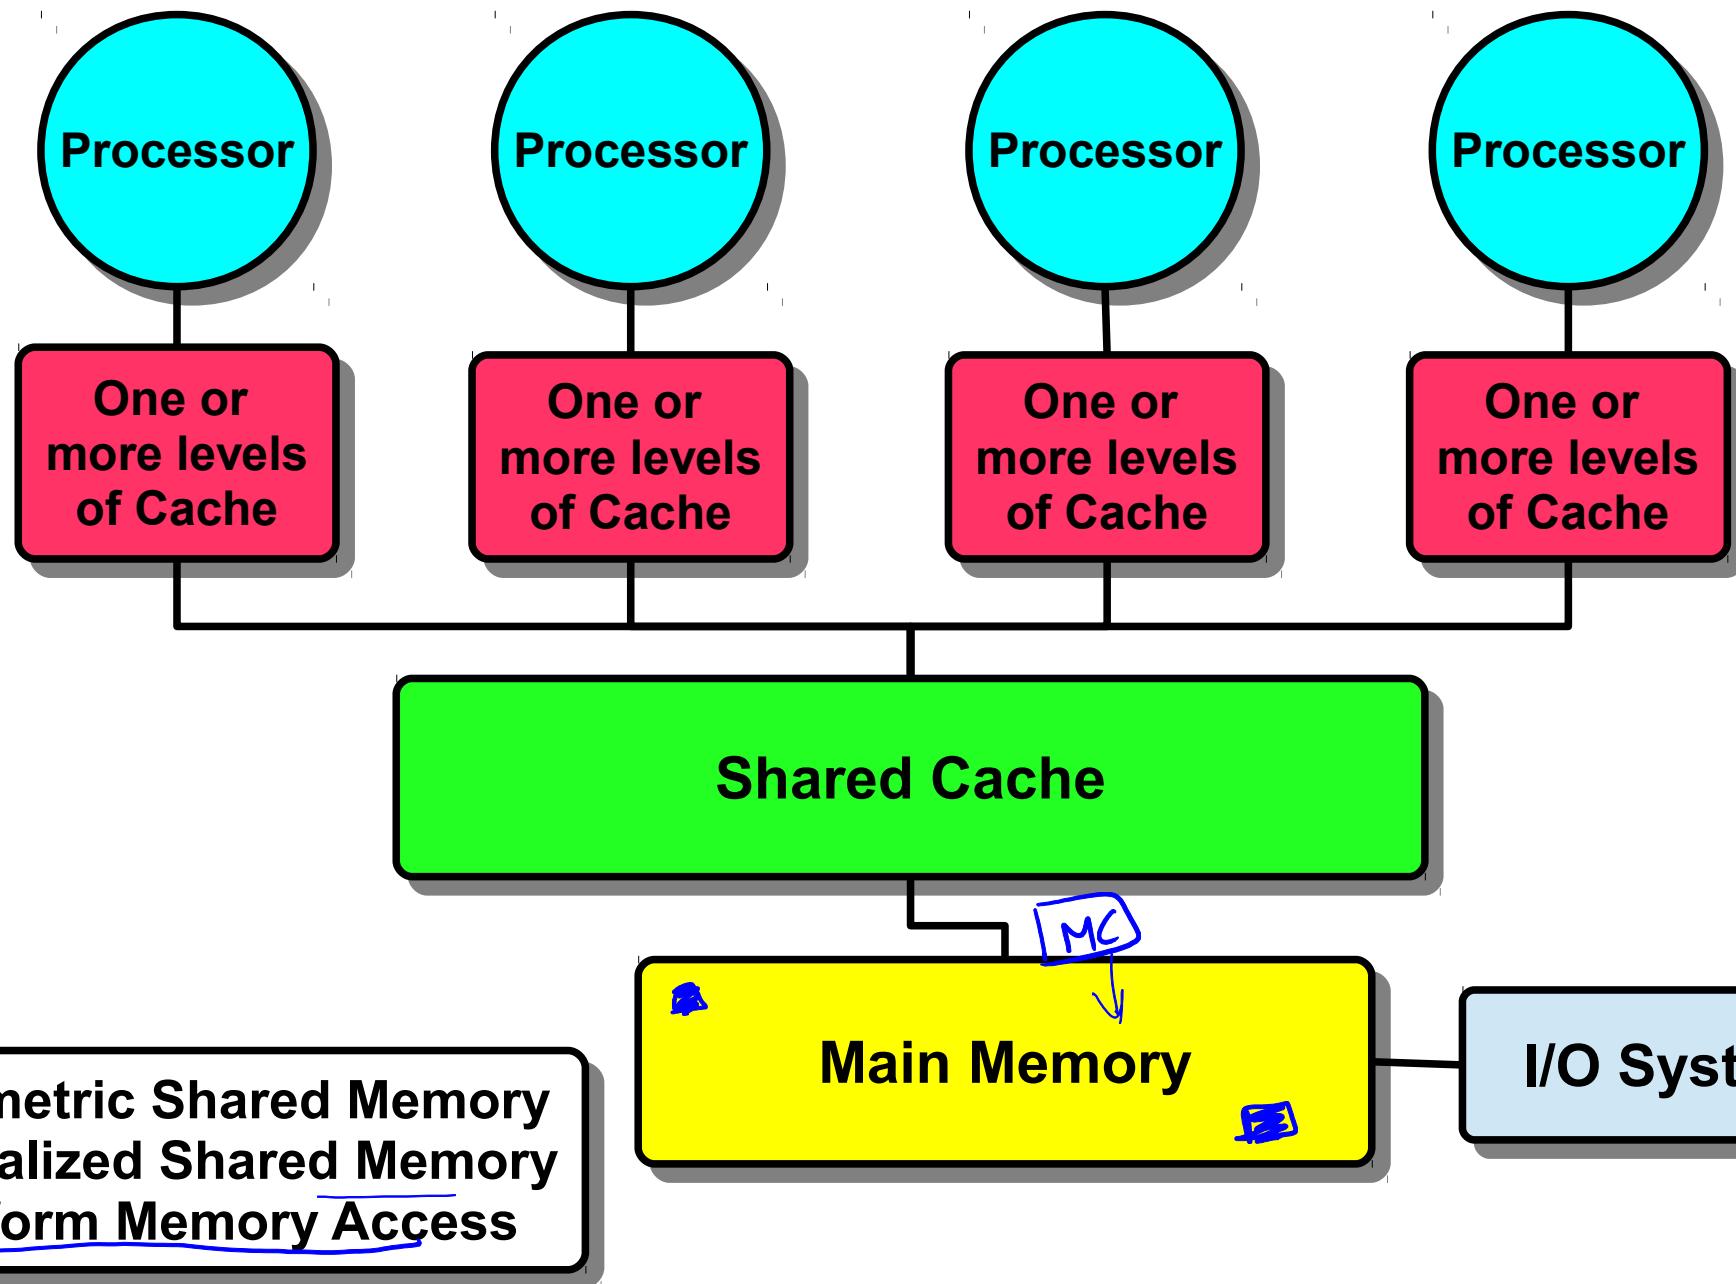
Symmetric Multiprocessor (SMP)



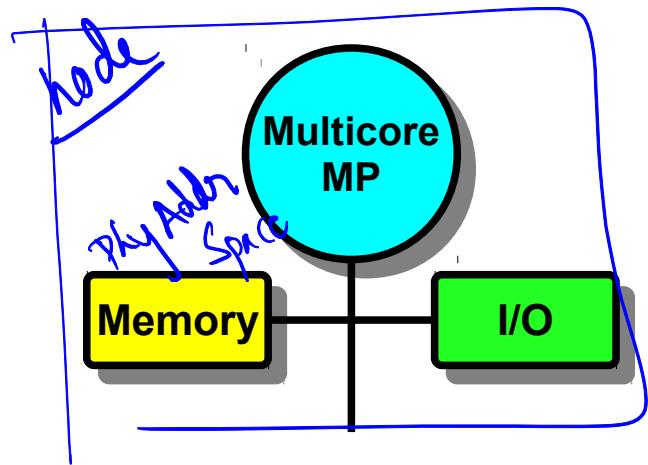
Symmetric Multiprocessor (SMP)



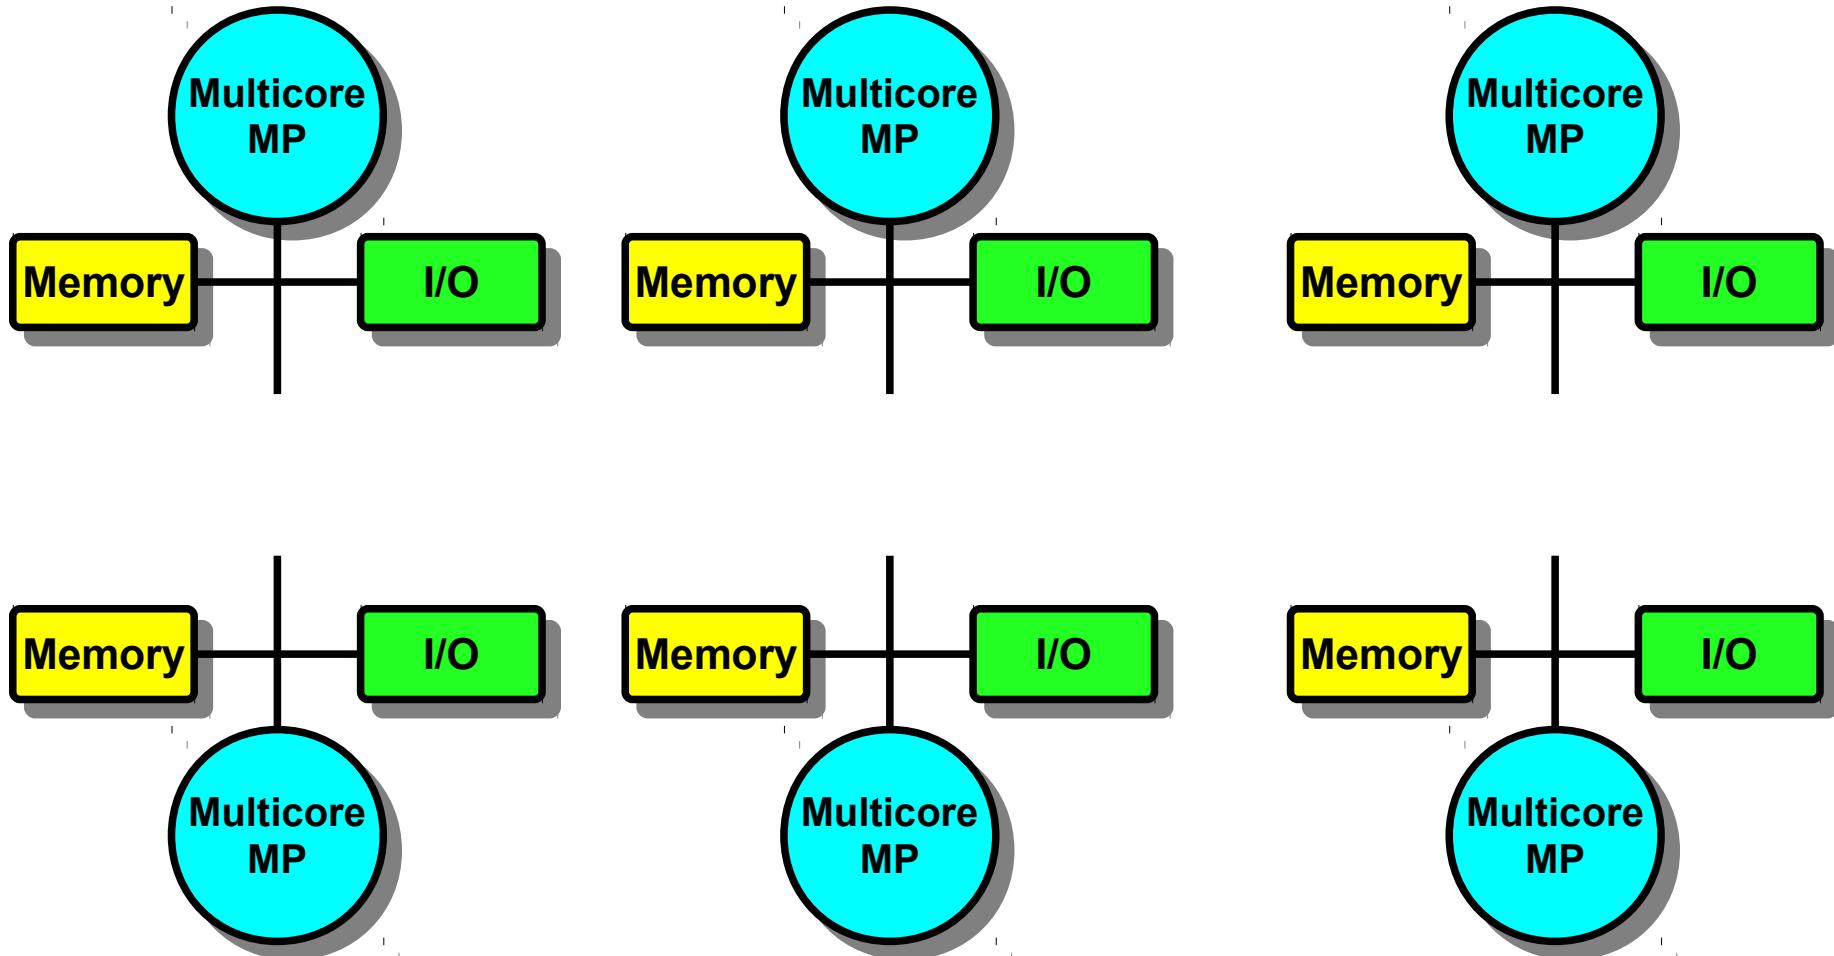
Symmetric Multiprocessor (SMP)



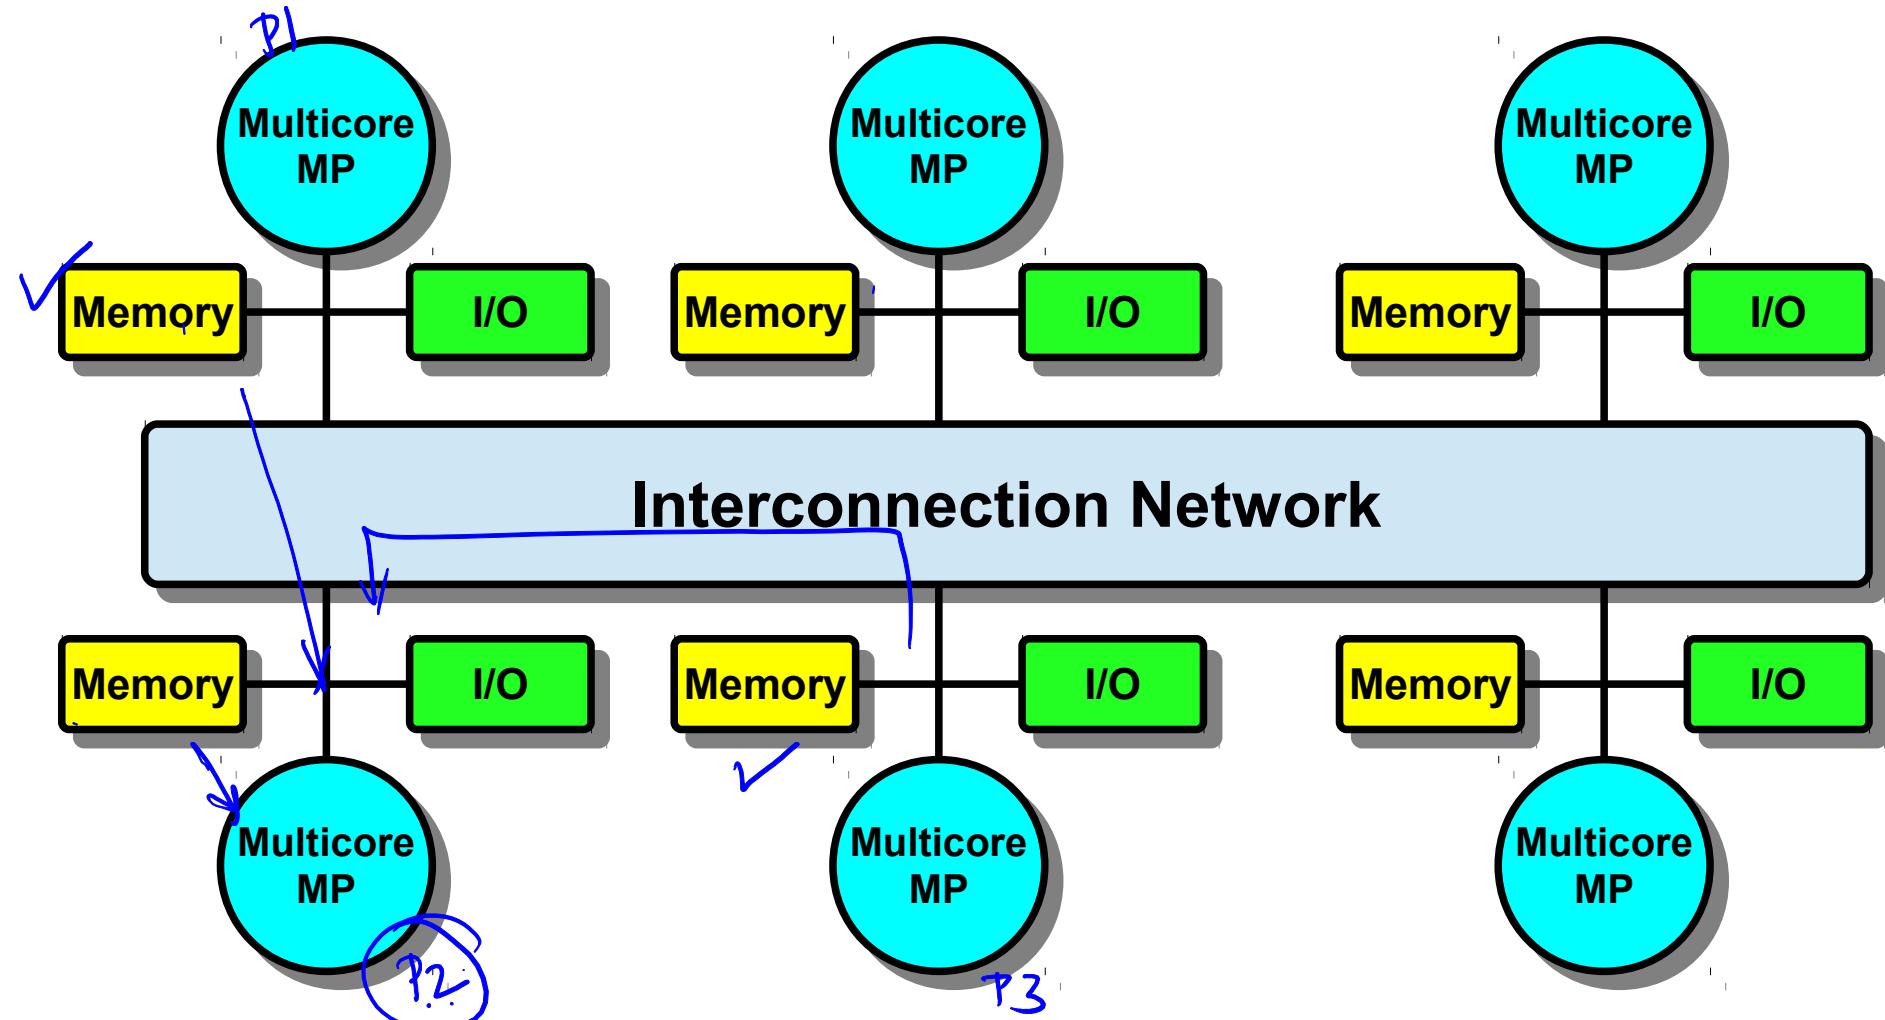
Distributed Shared Memory



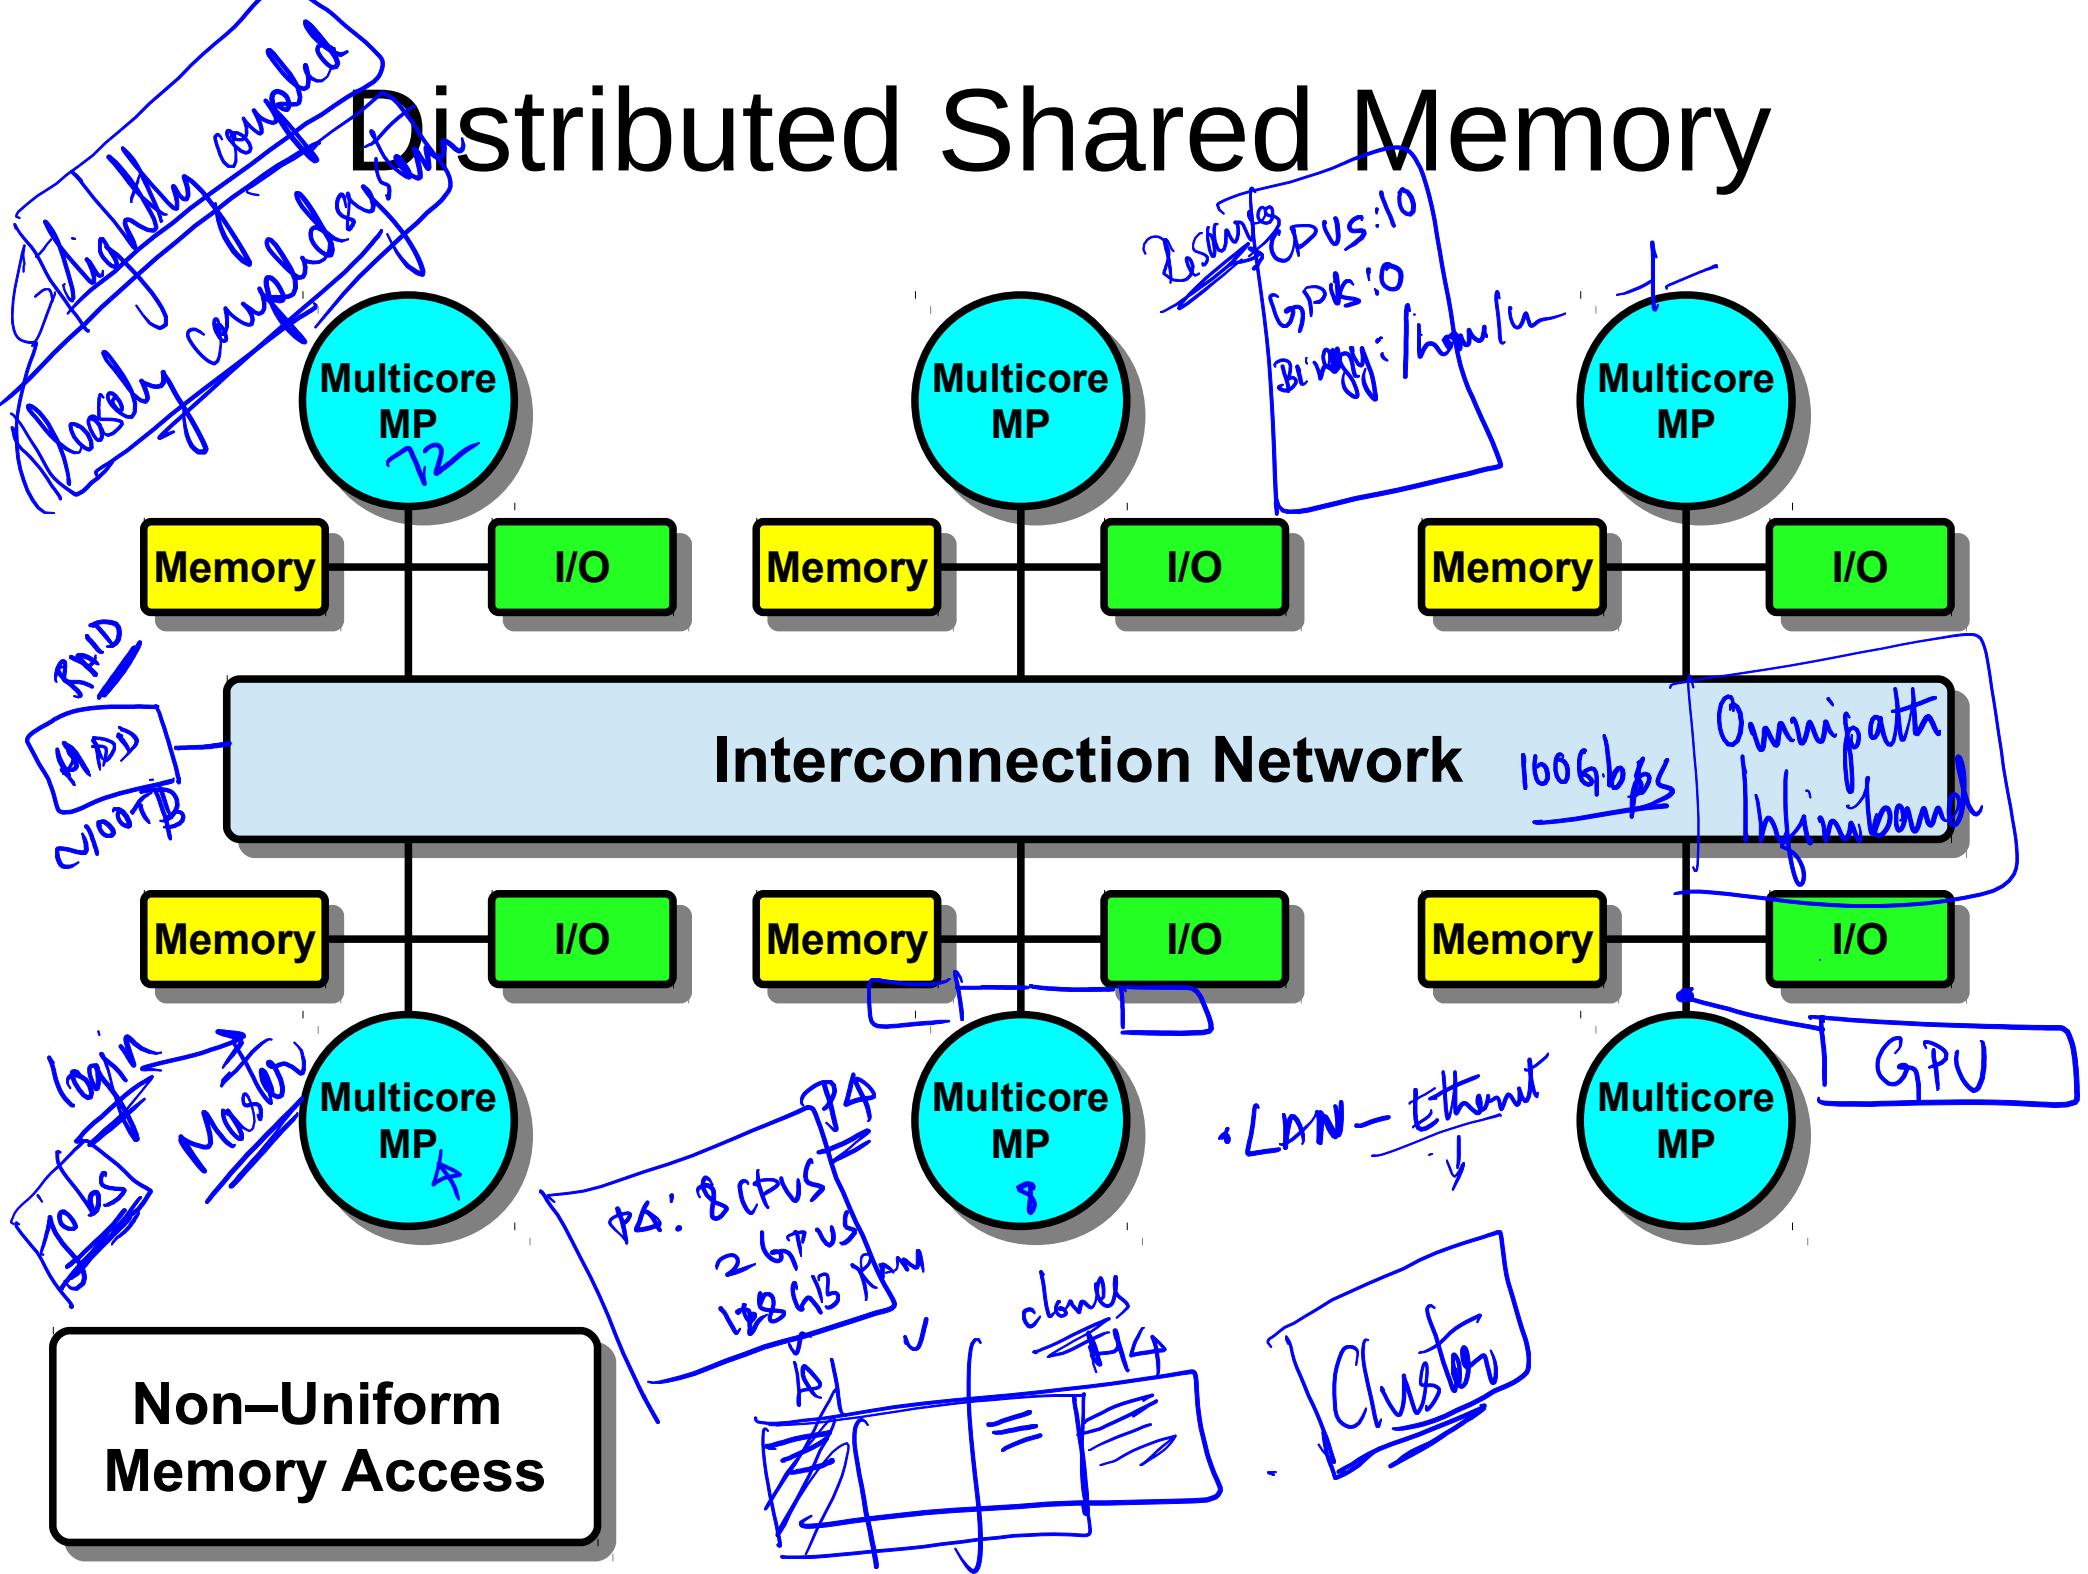
Distributed Shared Memory



Distributed Shared Memory



Distributed Shared Memory



Distributed Shared Memory

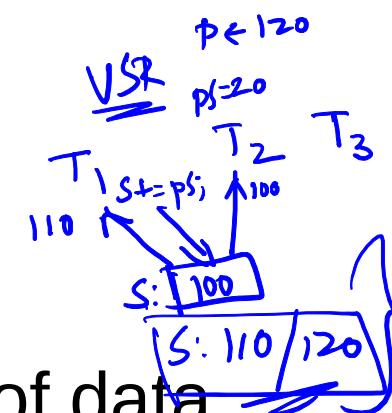
- Scalability
- High memory bandwidth demands
- Low memory access latency to local memory
- Communication infrastructure is complex

Shared Memory vs. Message Passing

- **Shared Memory Machine:** processors share the same physical address space
 -
 -
- **Message Passing Machine:** Memory is private
 -
 -

Shared Memory vs. Message Passing

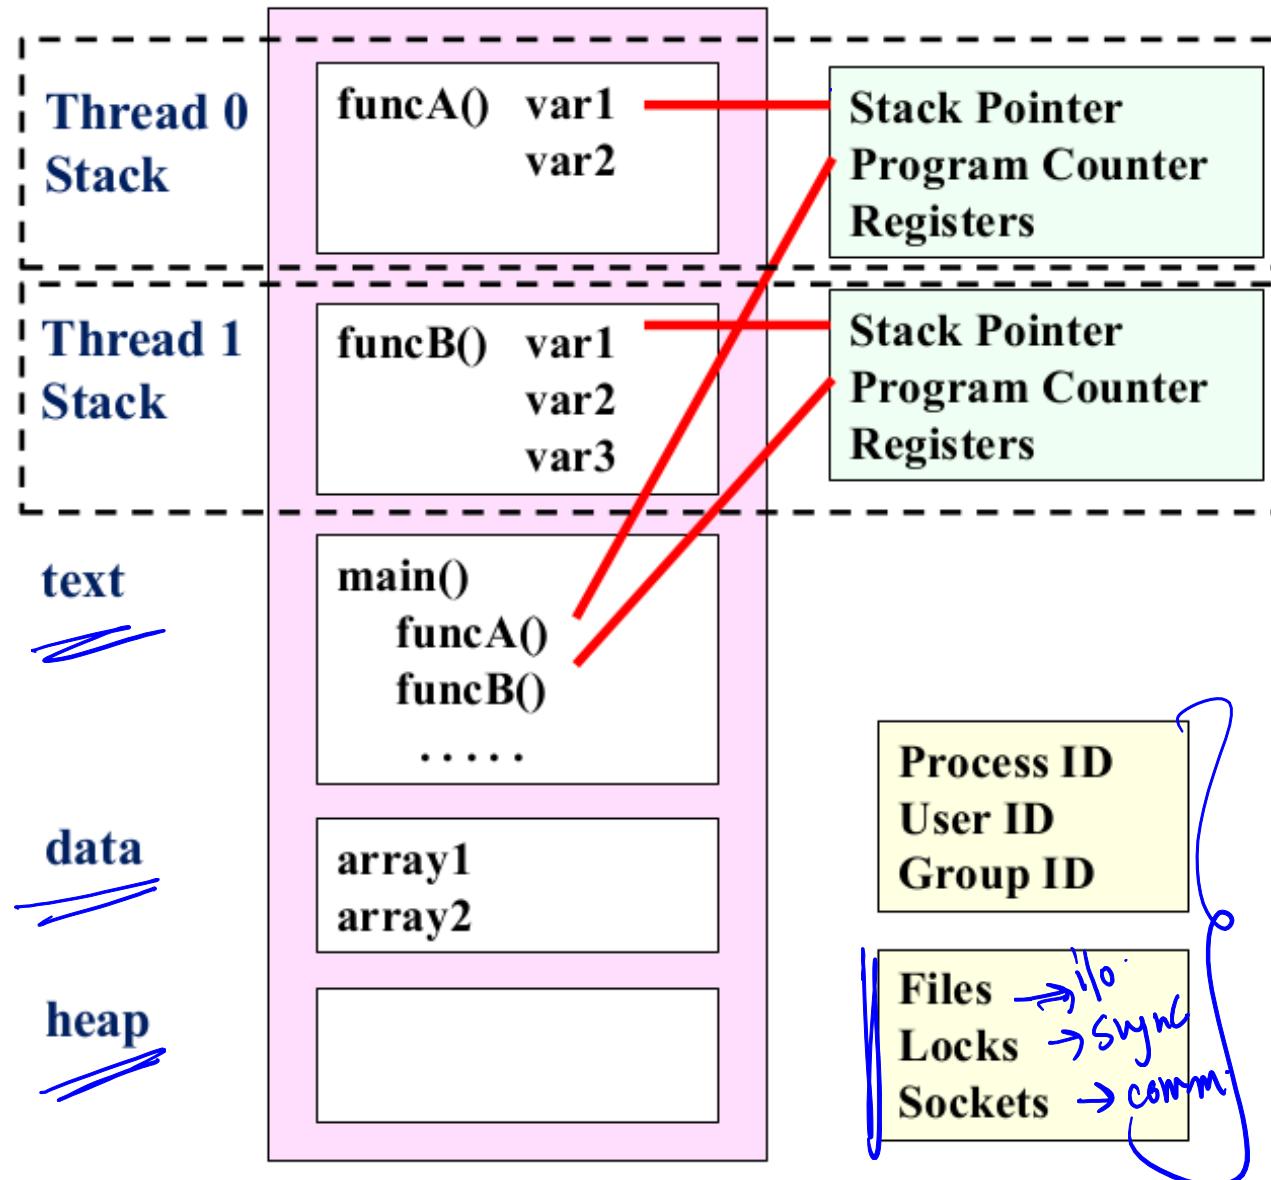
- **Shared Memory Machine:** processors share the same physical address space
 - Implicit communication via loads and stores
 - Needs locks, fences and critical sections to synchronize access
 - No need to know Destination on generation of data
(can store in memory and user of data can pick up later)
 - Easy for multiple threads accessing a shared table
 - Eg. OpenMP
- **Message Passing Machine:** Memory is private



Shared Memory vs. Message Passing

- **Shared Memory Machine:** processors share the same physical address space
- **Message Passing Machine:** Memory is private
 - Explicit send/receive to communicate
 - Message contains data and synchronization
 - Need to know Destination on generation of data (send)
 - Easy for Producer-Consumer
 - Eg. MPI

Threads vs. Processes

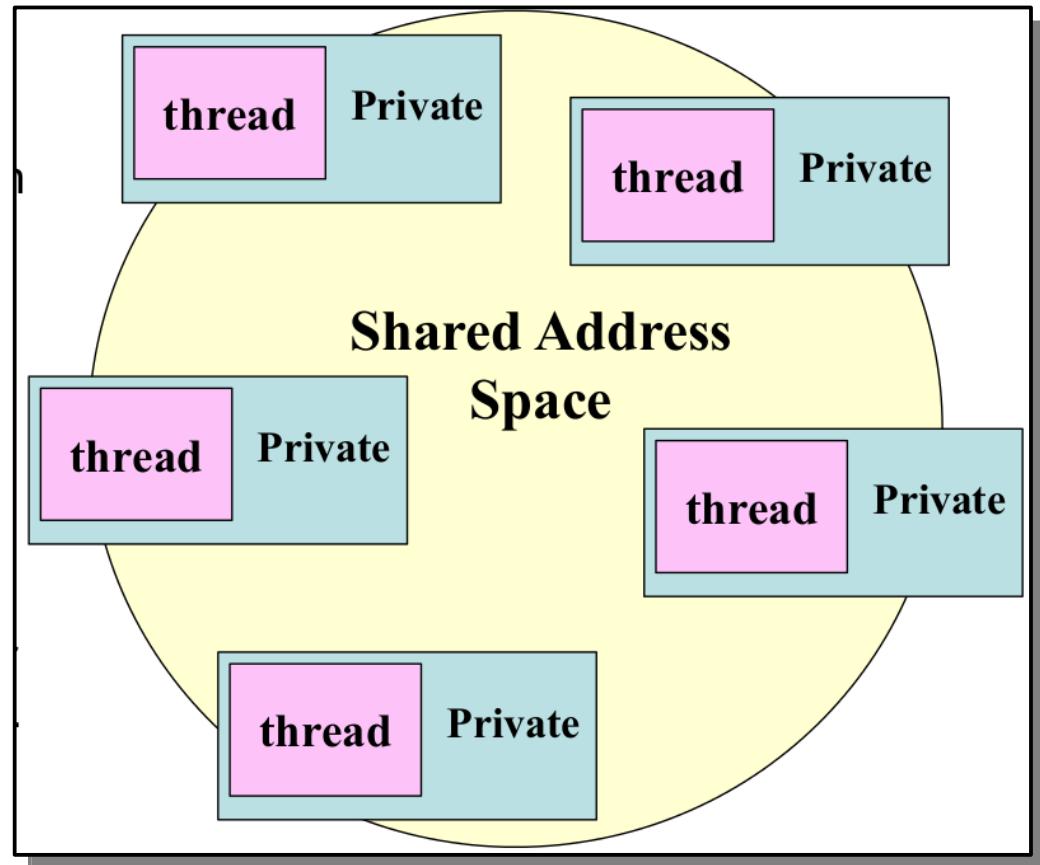


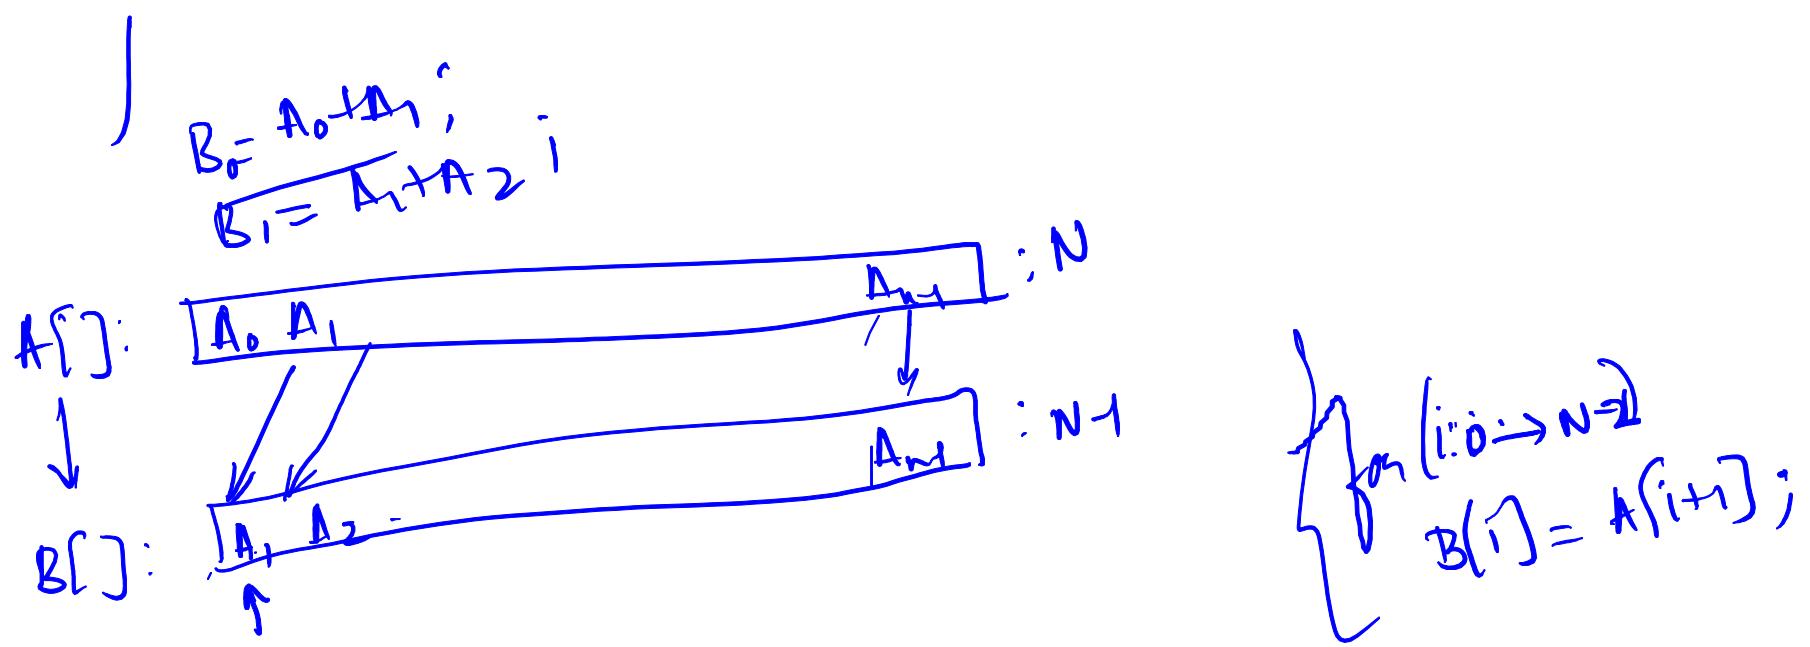
Threads:

- Threads are "light weight processes"
- Threads share Process state among multiple threads ... this greatly reduces the cost of switching context.

Shared Memory Program

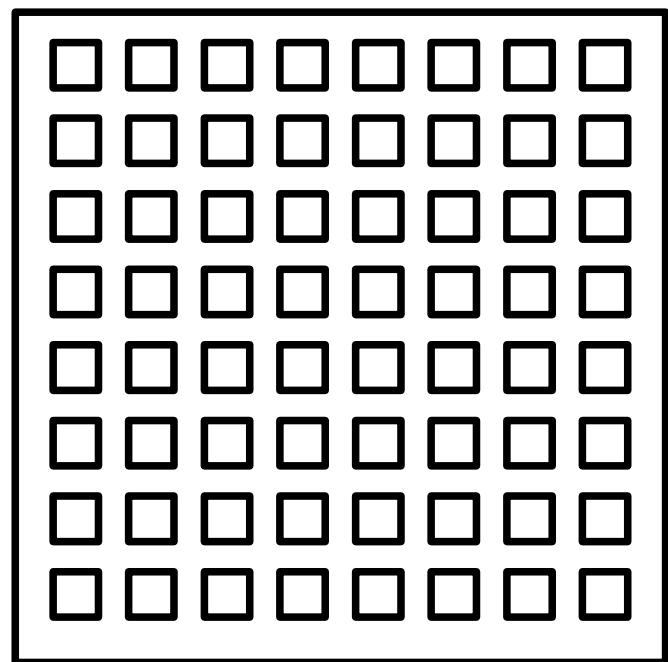
- One process, multiple threads.
- Threads interact through reads/writes to a shared address space.
- OS scheduler decides when to run which threads.
Interleaved for fairness.
- Synchronization to assure every legal order results in correct results.



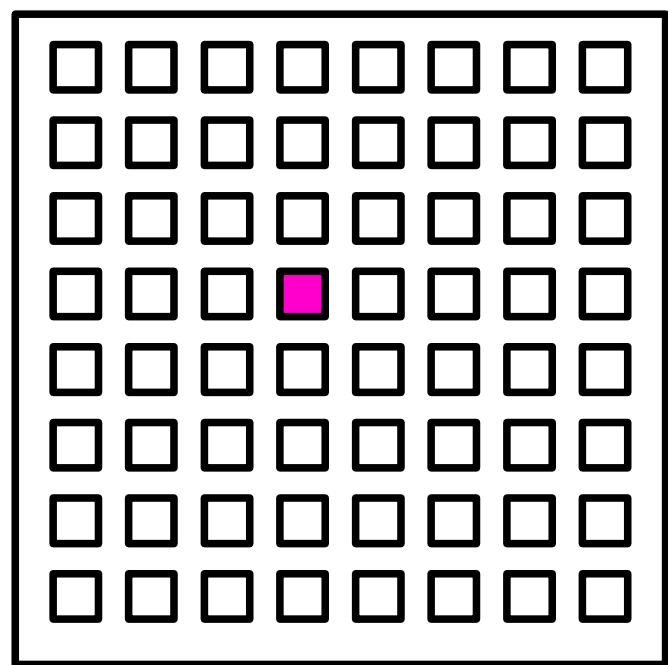


tasks: $N-k$
 index: $\dots : N-1$ $\overset{\text{tid}}{\text{tid}}$
threads: $N-1 : 0, 1, \dots, \overset{N-2}{\text{tid}}$
task for thread: update 1 element in B.
, $B[\text{tid}] = A[\text{tid}+1];$

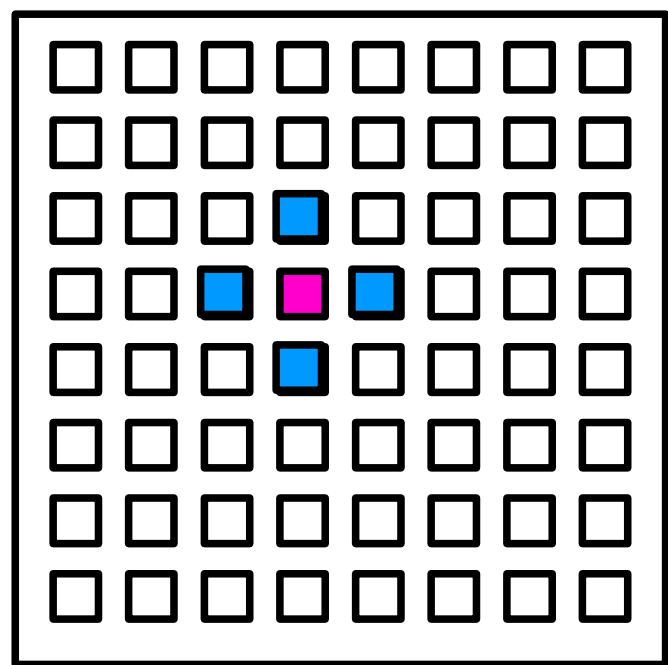
Ocean Kernel



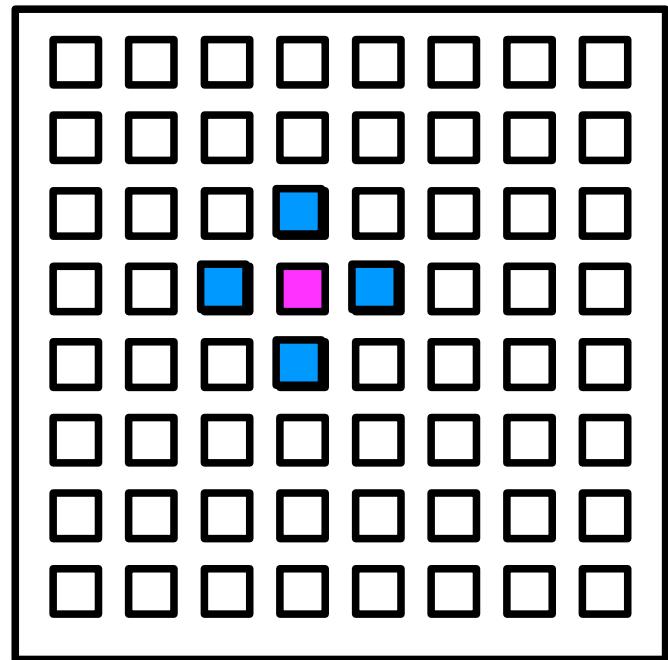
Ocean Kernel



Ocean Kernel

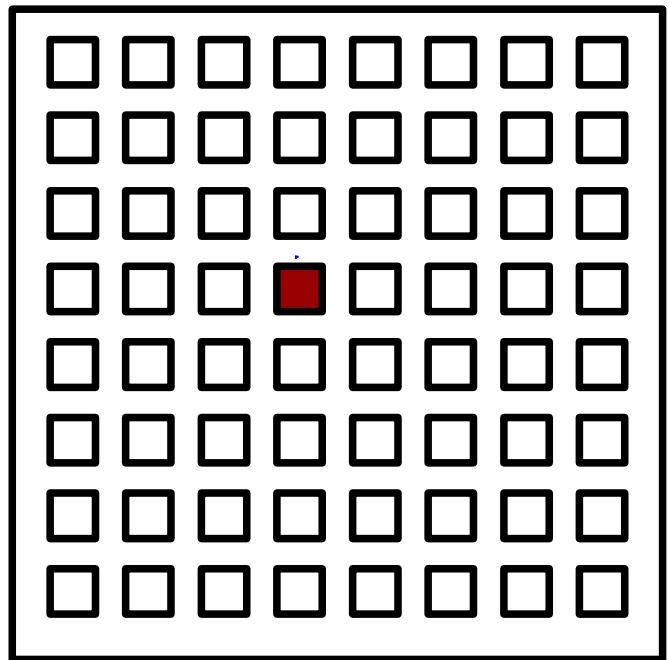


Ocean Kernel



$A[i,j] \leftarrow 0.2 * (A[i,j] + \text{neighbors});$ ↙

Ocean Kernel


$$A[i,j] \leftarrow 0.2 * (A[i,j] + \text{neighbors});$$

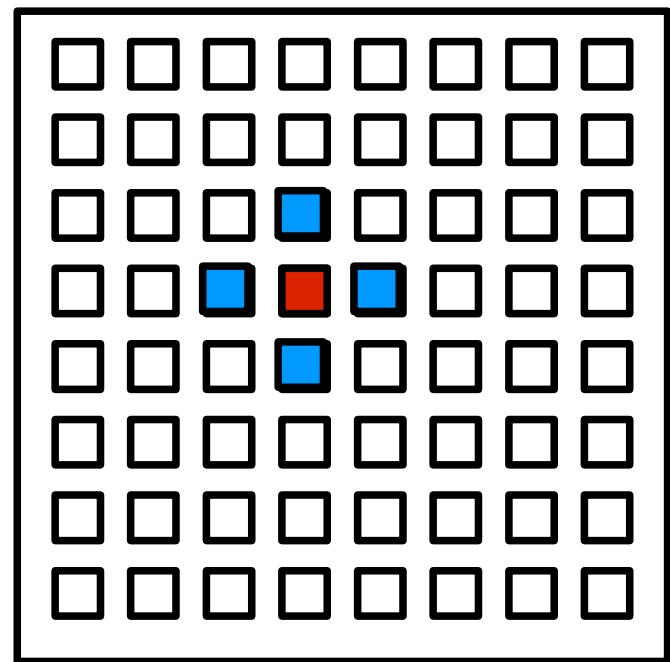
Ocean Kernel

Procedure Solve(A)

begin

```
{ for i < 1 to n do  
    for j < 1 to n do  
        temp = A[i,j];  
        new → A[i,j] ← 0.2 * (A[i,j] + neighbors);  
        diff += abs(A[i,j] – temp);  
    end for  
    ↑ new ↑ old  
end for
```

end procedure



Ocean Kernel

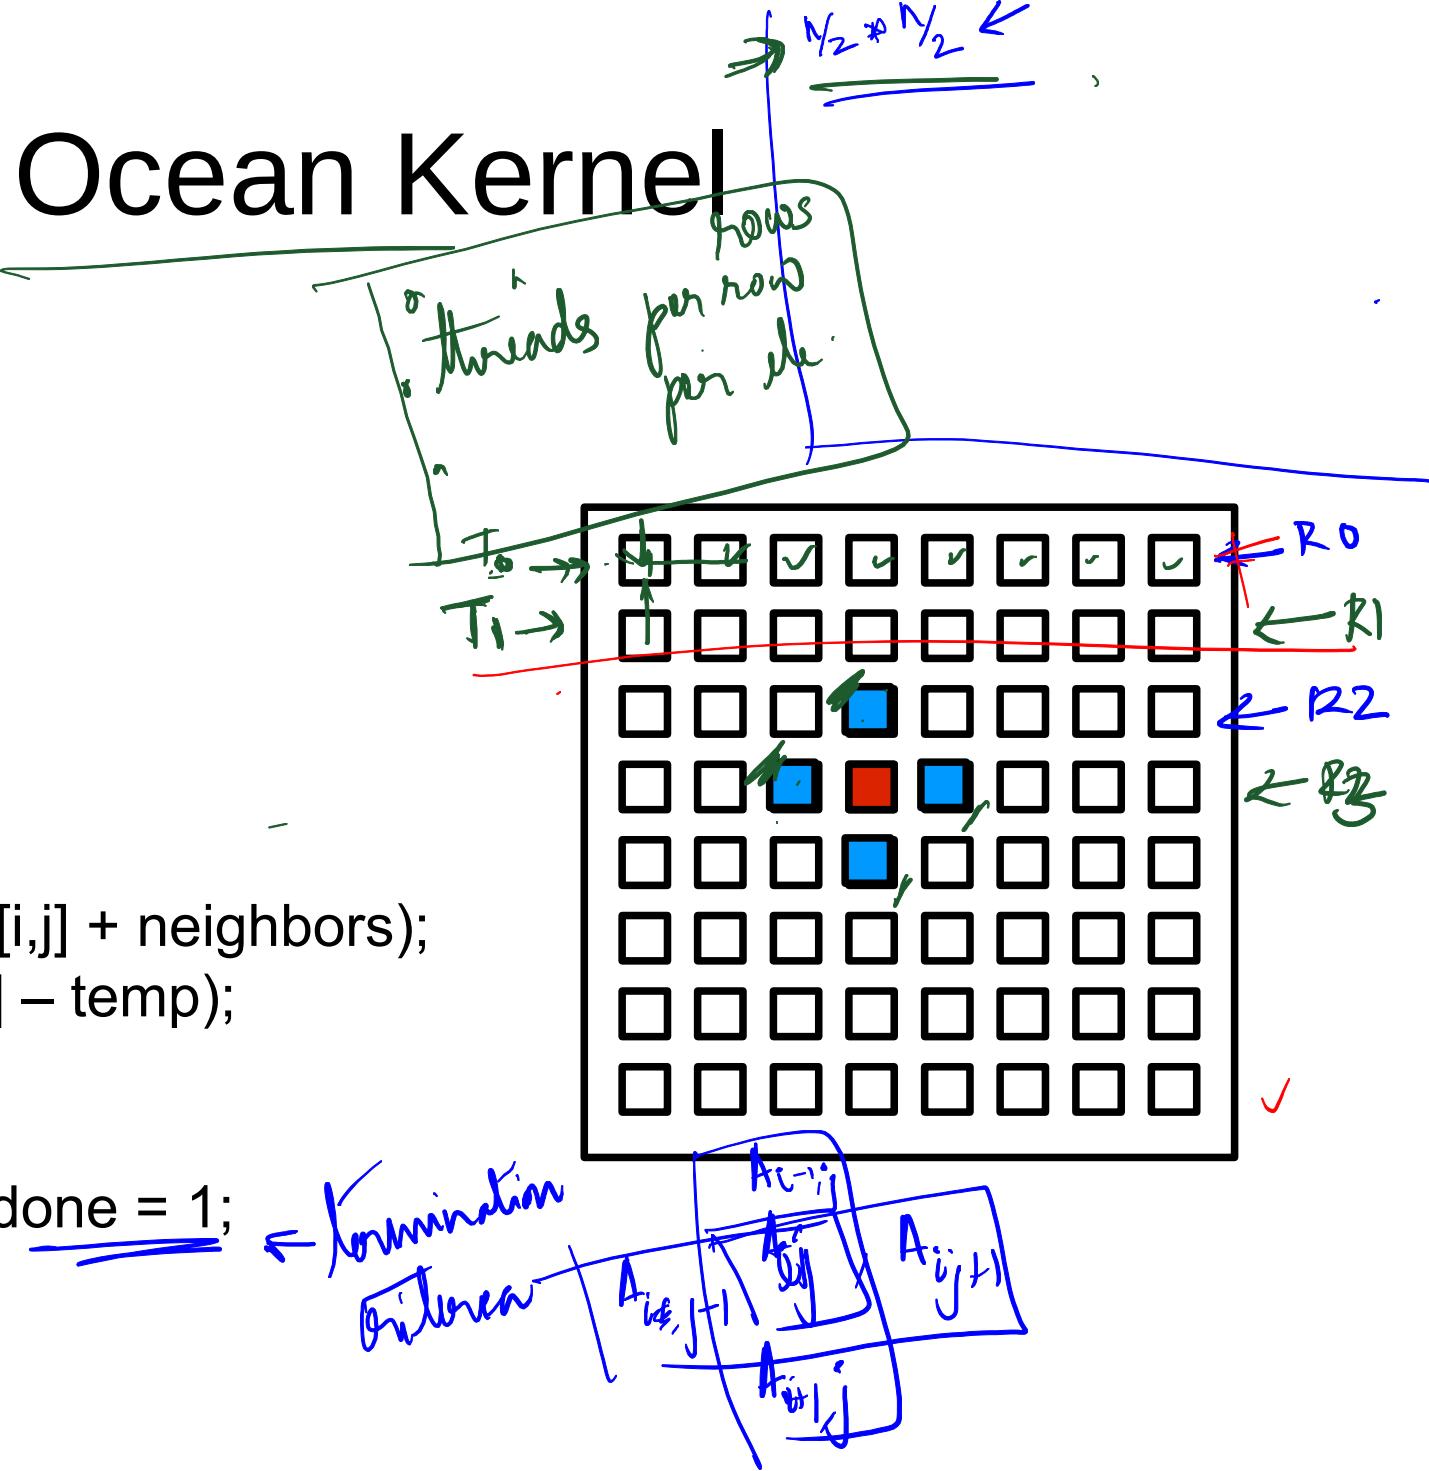
Procedure Solve(A)

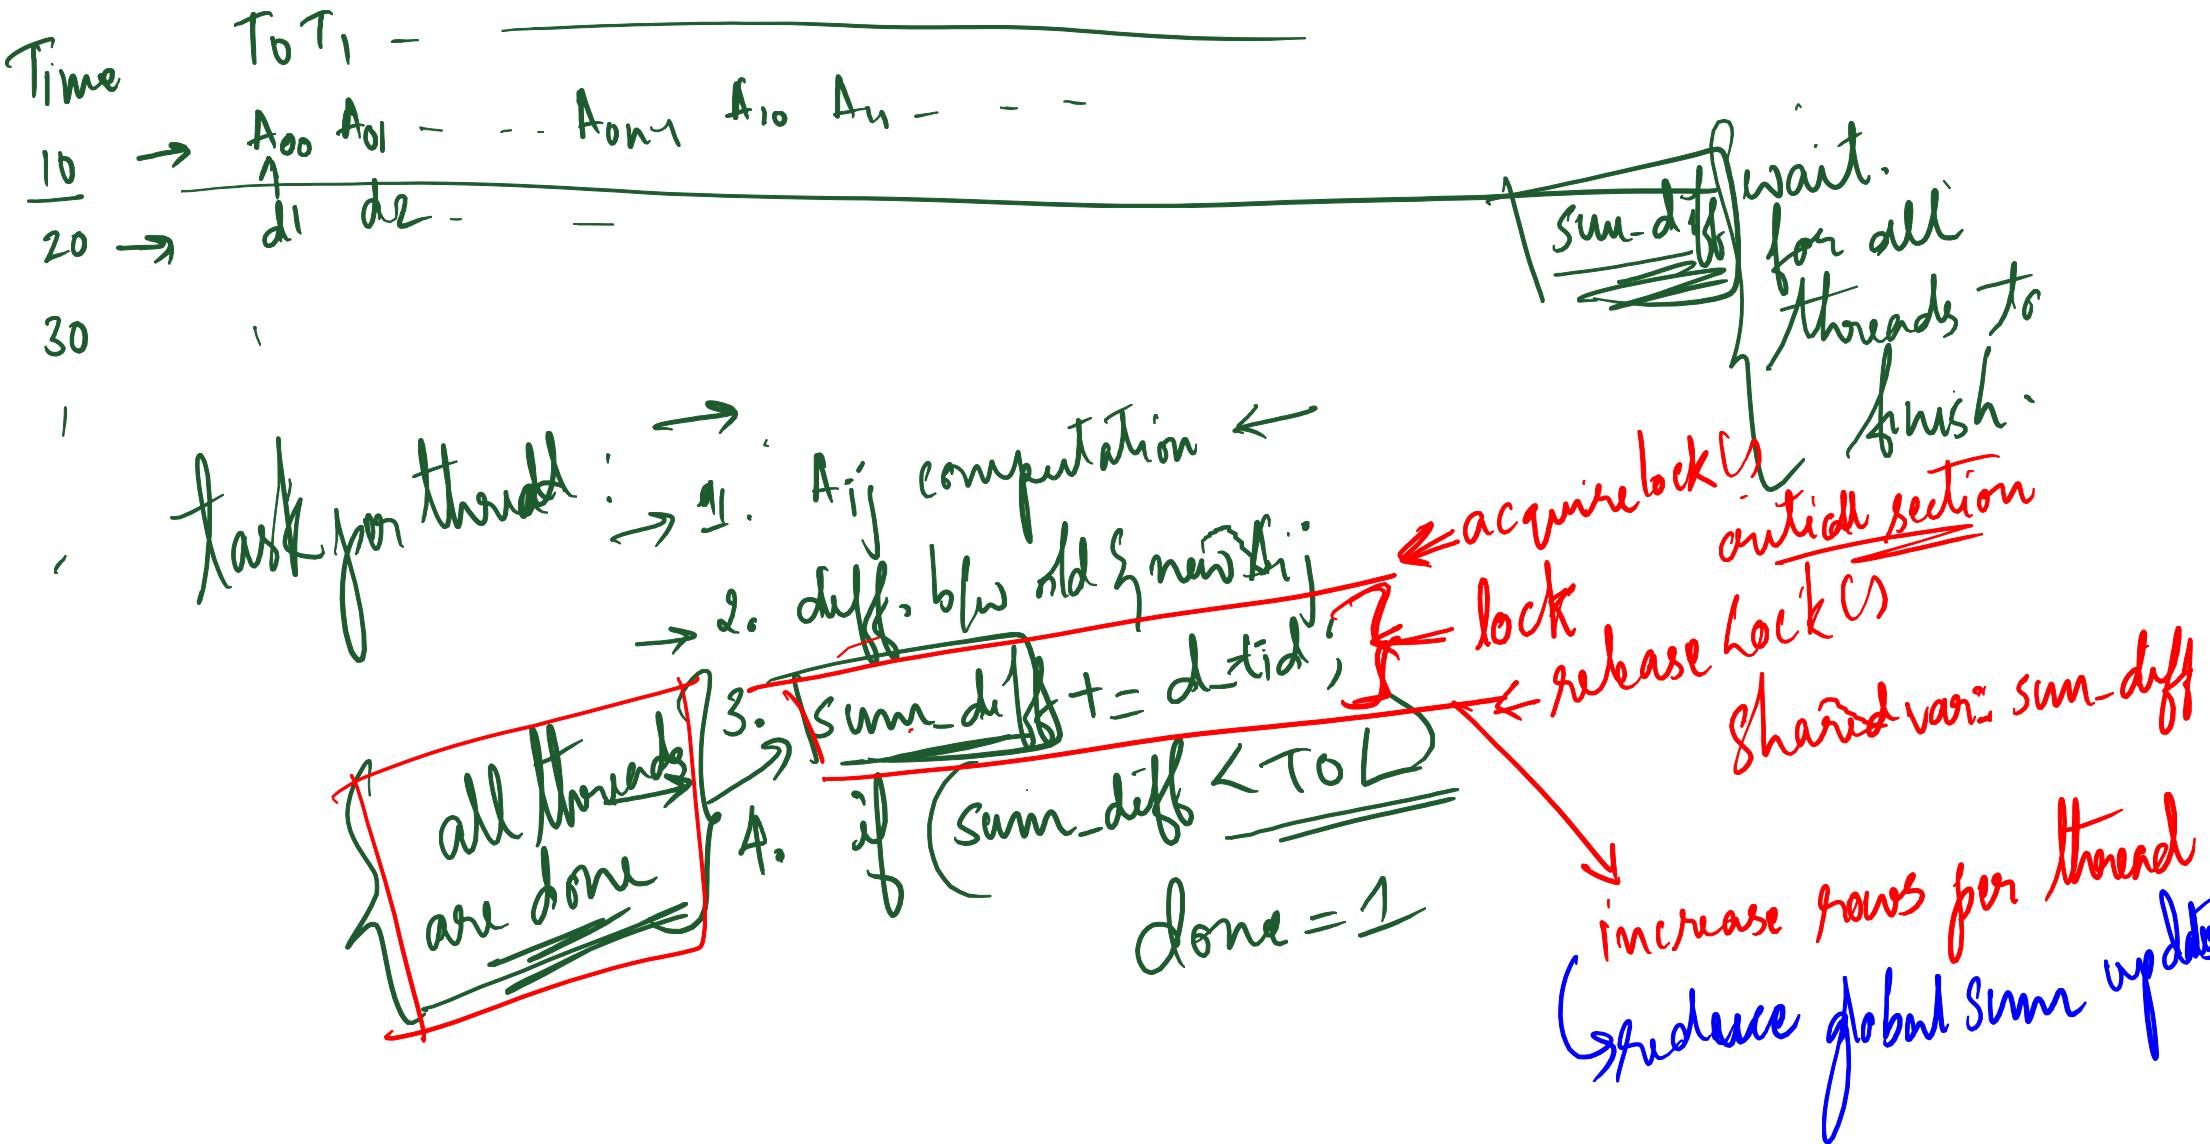
begin

```

✓while (!done) do
    diff = 0;
    for i ← 1 to n do
        for j ← 1 to n do
            temp = A[i,j];
            A[i,j] ← 0.2 * (A[i,j] + neighbors);
            diff += abs(A[i,j] – temp);
        end for
    end for
    if (diff < TOL) then done = 1;
end while
0.01
end procedure

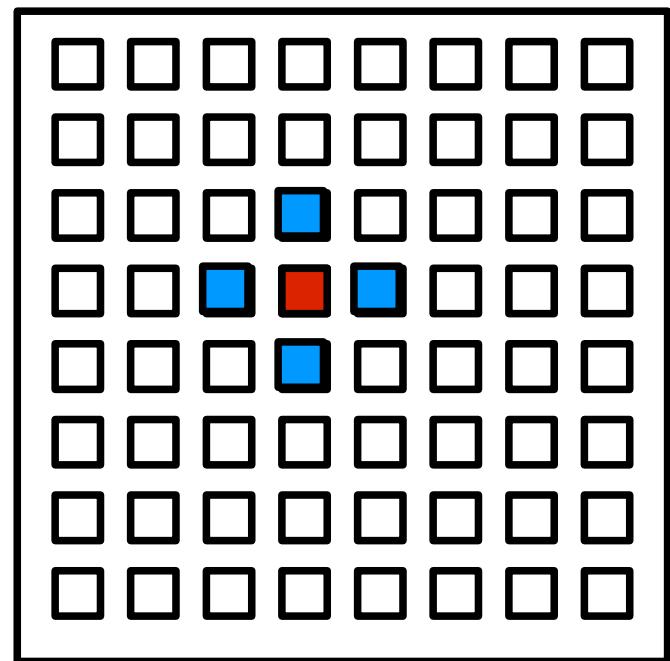
```





Ocean Kernel

```
Procedure Solve(A)
begin
    diff = done = 0;
    while (!done) do
        diff = 0;
        for i  $\leftarrow$  1 to n do
            for j  $\leftarrow$  1 to n do
                temp = A[i,j];
                A[i,j]  $\leftarrow$  0.2 * (A[i,j] + neighbors);
                diff += abs(A[i,j] – temp);
            end for
        end for
        if (diff < TOL) then done = 1;
    end while
end procedure
```

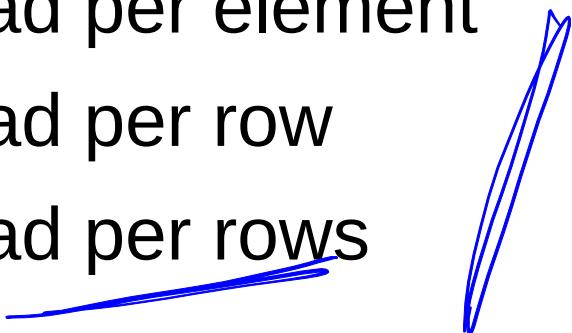


Shared Address – Ocean Kernel

- List the tasks.
- Create how many threads?
-

Shared Address – Ocean Kernel

- List the tasks.
- Create how many threads?
 - Thread per element
 - Thread per row
 - Thread per rows



Shared Address – Ocean Kernel

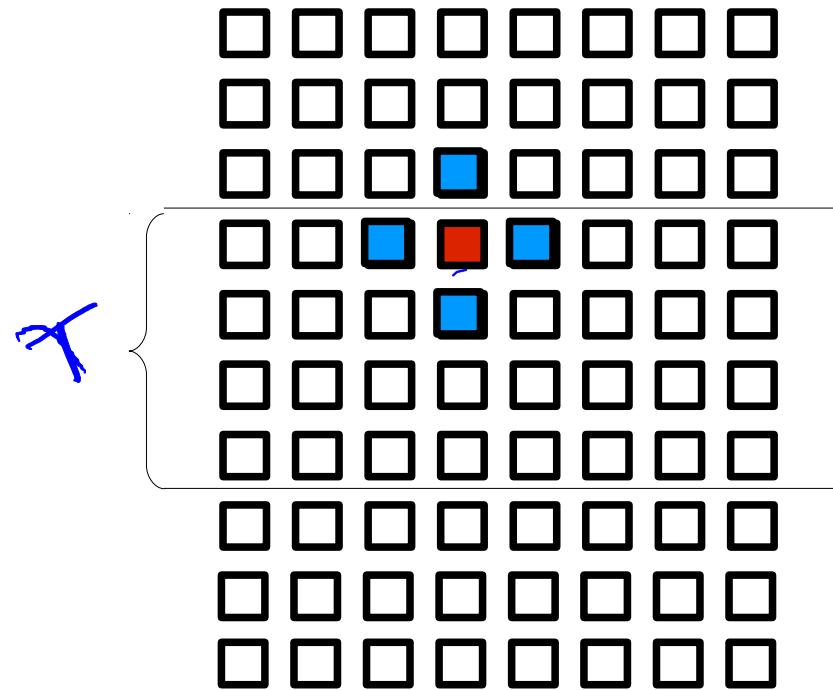
- List the tasks.
- Create how many threads?
- Tasks per thread?

Ocean Kernel – Threads

- One thread for a few rows

Ocean Kernel – Threads

- One thread for a few rows



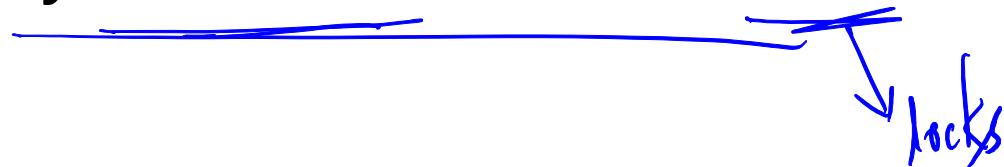
new diff $A_{ij} \leftarrow \text{avg} (A_{ij-\text{old}} + 4 \text{ neighbours})$
 ≤ 0.2

Ocean Kernel – Threads

- One thread for a few rows
- Thread must update the partial sum (diff) after averaging every element in its rows

Ocean Kernel – Threads

- One thread for a few rows
- Thread must update the partial sum (diff)
 - Synchronization, Mutual exclusion



Ocean Kernel – Threads

- One thread for a few rows
- Thread must update the partial sum (diff)
 - Synchronization, Mutual exclusion
 - Lock

Ocean Kernel – Threads

- One thread for a few rows
- Thread must update the partial sum (diff)
 - Synchronization, Mutual exclusion
 - Lock
 - Acquire, Release



Ocean Kernel – Threads

- One thread for a few rows
- Thread must update the partial sum (diff)
- After all threads update diff, check if exceeds threshold

Ocean Kernel – Threads

- One thread for a few rows
- Thread must update the partial sum (diff)
- After all threads update diff, check if exceeds threshold
 - Barrier

Ocean Kernel – Threads

- One thread for a few rows
- Thread must update the partial sum (diff)
- After all threads update diff, check if exceeds threshold
 - Barrier
 - All threads call barrier
 - Every thread waits till peer reaches barrier call

Ocean Kernel – Threads

- One thread for a few rows
- Thread must update the partial sum (diff)
- After all threads update diff, check if exceeds threshold
 - Yes – terminate thread
 - No – repeat iteration

Shared Address – Ocean Kernel

Each thread executes instructions
from a procedure – Solve(A)

procedure Solve(A)

Shared Address – Ocean Kernel

procedure Solve(A)

Thread id

Thread identifies its start row and end row

Shared Address – Ocean Kernel

Thread identifies its start row and end row

procedure Solve(A)

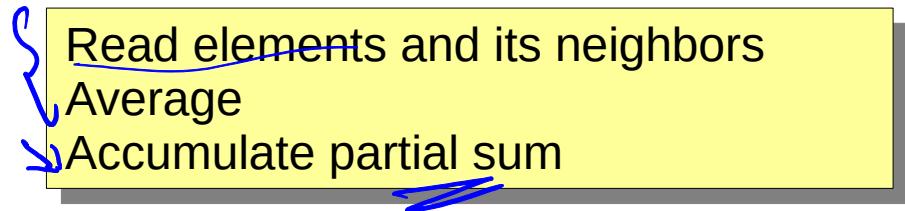
int mymin = 1 + (pid * n/procs);
int mymax = mymin + n/nprocs - 1;

102 400
||

Shared Address – Ocean Kernel

```
procedure Solve(A)
```

```
    int mymin = 1 + (pid * n/procs);  
    int mymax = mymin + n/nprocs -1;
```



Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);  
int mymax = mymin + n/nprocs -1;
```

```
→ for i ← mymin to mymax  
      for j ← 1 to n do  
          /* avg neighbours */  
          /* accumulate mydiff */  
      endfor  
  endfor  
→ pa[i][j] = ...
```

first row *last row*

Read elements and its neighbours
Average
Accumulate partial sum

Shared Address – Ocean Kernel

```
procedure Solve(A)
```

```
    int mymin = 1 + (pid * n/procs);  
    int mymax = mymin + n/nprocs -1;
```

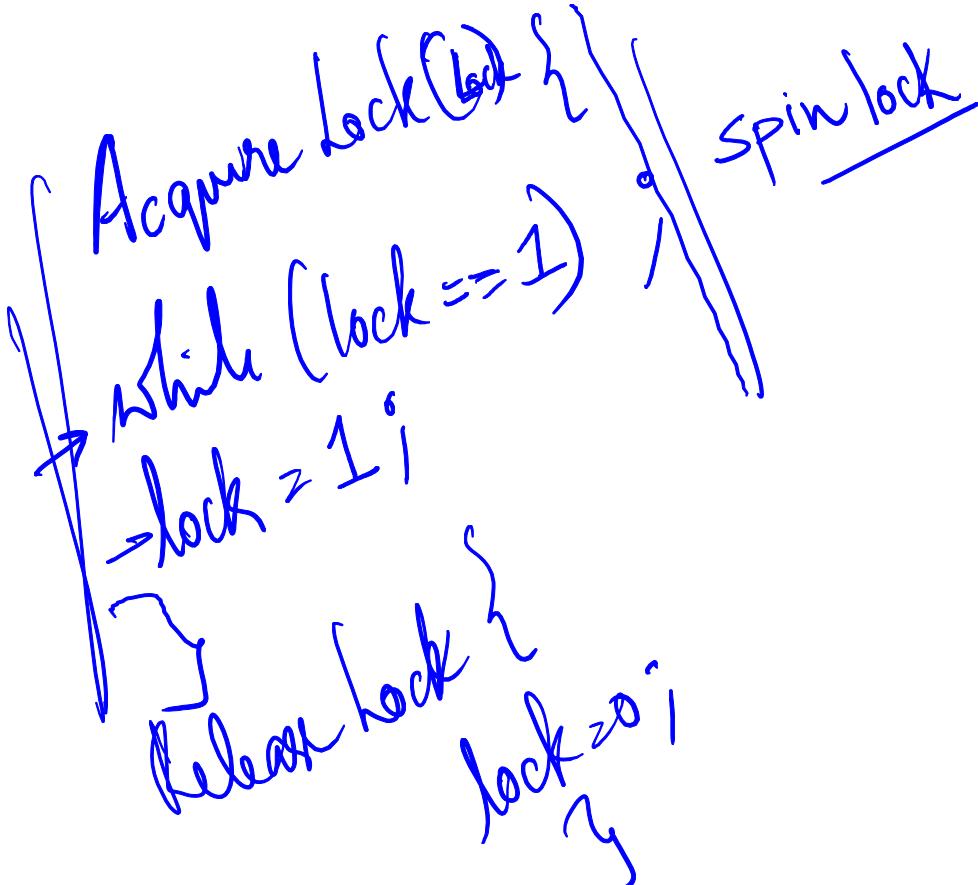
```
    for i ← mymin to mymax  
        for j ← 1 to n do  
            /* avg neighbours */  
            /* accumulate mydiff */  
        endfor  
    endfor
```

— —

Update diff

global

Shared Address – Ocean Kernel



procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);  
int mymax = mymin + n/nprocs - 1;
```

```
for i ← mymin to mymax  
  for j ← 1 to n do  
    /* avg neighbours */  
    /* accumulate mydiff */  
  endfor  
endfor
```

```
LOCK(diff_lock);  
diff += mydiff;  
UNLOCK(diff_lock);
```

Shared Address – Ocean Kernel

```
procedure Solve(A)
```

```
    int mymin = 1 + (pid * n/procs);  
    int mymax = mymin + n/nprocs -1;
```

```
    for i ← mymin to mymax  
        for j ← 1 to n do  
            /* avg neighbours */  
            /* accumulate mydiff */  
        endfor  
    endfor  
    LOCK(diff_lock);  
    diff += mydiff;  
    UNLOCK(diff_lock);
```

Wait for all threads to update diff

Shared Address – Ocean Kernel

```
procedure Solve(A)
```

```
    int mymin = 1 + (pid * n/procs);
    int mymax = mymin + n/nprocs -1;
```

```
    for i ← mymin to mymax
        for j ← 1 to n do
            /* avg neighbours */
            /* accumulate mydiff */
        endfor
    endfor
    LOCK(diff_lock);
    diff += mydiff;
    UNLOCK(diff_lock);
    BARRIER (bar1, nprocs);
```

Wait for all threads to update diff

Shared Address – Ocean Kernel

```
procedure Solve(A)
```

```
    int mymin = 1 + (pid * n/procs);  
    int mymax = mymin + n/nprocs -1;
```

```
    for i ← mymin to mymax  
        for j ← 1 to n do  
            /* avg neighbours */  
            /* accumulate mydiff */  
        endfor  
    endfor  
    LOCK(diff_lock);  
    diff += mydiff;  
    UNLOCK(diff_lock);  
    BARRIER (bar1, nprocs);
```

globaldiff is now ready.

Check for threshold
Terminate or Repeat

vs. globaldiff

Shared Address – Ocean Kernel

```
procedure Solve(A)
```

```
    int mymin = 1 + (pid * n/procs);  
    int mymax = mymin + n/nprocs -1;
```

```
    for i ← mymin to mymax  
        for j ← 1 to n do  
            /* avg neighbours */  
            /* accumulate mydiff */  
        endfor  
    endfor  
    LOCK(diff_lock);  
    diff += mydiff;  
    UNLOCK(diff_lock);  
    BARRIER (bar1, nprocs);  
    if (diff < TOL) then done = 1;
```

Check for threshold
Terminate or Repeat

Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);  
int mymax = mymin + n/nprocs - 1;
```

```
while (!done) do
```

```
    mydiff = diff = 0;
```

```
    for i ← mymin to mymax
```

```
        for j ← 1 to n do
```

```
            /* avg neighbours */
```

```
            /* accumulate mydiff */
```

```
        endfor
```

```
    endfor
```

```
    LOCK(diff_lock);
```

```
    diff += mydiff;
```

```
    UNLOCK(diff_lock);
```

```
    BARRIER (bar1, nprocs);
```

```
    if (diff < TOL) then done = 1;
```

```
endwhile
```

fast T₂

Check for threshold
Terminate or Repeat

slow T₁

wait

Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);  
int mymax = mymin + n/nprocs - 1;
```

```
while (!done) do
```

```
    mydiff = diff = 0; ;
```

```
    for i ← mymin to mymax
```

```
        for j ← 1 to n do
```

```
            /* avg neighbours */
```

```
            /* accumulate mydiff */
```

```
        endfor
```

```
    endfor
```

```
    LOCK(diff_lock);
```

```
    diff += mydiff;
```

```
    UNLOCK(diff_lock);
```

```
    BARRIER (bar1, nprocs);
```

```
    if (diff < TOL) then done = 1;
```

```
    BARRIER()
```

```
endwhile
```

What happens if one thread runs ahead and executes

Slow T₂

T₆ — T₁

for i ← mymin to mymax

for j ← 1 to n do

/ avg neighbours */*

/ accumulate mydiff */*

endfor

endfor

LOCK(diff_lock);

diff += mydiff;

UNLOCK(diff_lock);

BARRIER (bar1, nprocs);

if (diff < TOL) then done = 1;

BARRIER()

Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);
int mymax = mymin + n/nprocs - 1;
while (!done) do
    mydiff = diff = 0; <
```

```
    for i ← mymin to mymax
        for j ← 1 to n do
            /* avg neighbours */
            /* accumulate mydiff */
        endfor
    endfor
```

```
    LOCK(diff_lock);
    diff += mydiff;
    UNLOCK(diff_lock);
    BARRIER (bar1, nprocs);
    if (diff < TOL) then done = 1;
    BARRIER
endwhile
```

What happens if one thread runs ahead and executes while other threads are still executing

Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);
int mymax = mymin + n/nprocs - 1;
while (!done) do
    mydiff = diff = 0;
```

for i \leftarrow mymin to mymax
for j \leftarrow 1 to n do

/* avg neighbours */
/* accumulate mydiff */

endfor

endfor

LOCK(diff_lock);

diff += mydiff;

UNLOCK(diff_lock);

BARRIER (bar1, nprocs);

if (diff < TOL) then done = 1;

endwhile

What happens if one thread runs ahead and executes while other threads are still executing

All thread should have completed threshold checking before any thread moves to the next iteration

Slow thread

Fast updating diff
Thread

Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);
int mymax = mymin + n/nprocs -1;
while (!done) do
    mydiff = diff = 0;

    for i ← mymin to mymax
        for j ← 1 to n do
            /* avg neighbours */
            /* accumulate mydiff */
        endfor
    endfor
    LOCK(diff_lock);
    diff += mydiff;
    UNLOCK(diff_lock);
    BARRIER (bar1, nprocs);
    if (diff < TOL) then done = 1;
    BARRIER (bar1, nprocs);
endwhile
```

What happens if one thread runs ahead and executes while other threads are still executing

All thread should have completed threshold checking before any thread moves to the next iteration

Shared Address – Ocean Kernel

```
procedure Solve(A)
```

```
    int mymin = 1 + (pid * n/procs);  
    int mymax = mymin + n/nprocs -1;  
    while (!done) do  
        mydiff = diff = 0;
```

```
        for i ← mymin to mymax  
            for j ← 1 to n do  
                /* avg neighbours */  
                /* accumulate mydiff */  
            endfor  
        endfor  
        LOCK(diff_lock);  
        diff += mydiff;  
        UNLOCK(diff_lock);  
        BARRIER (bar1, nprocs);  
        if (diff < TOL) then done = 1;  
        BARRIER (bar1, nprocs);  
    endwhile
```

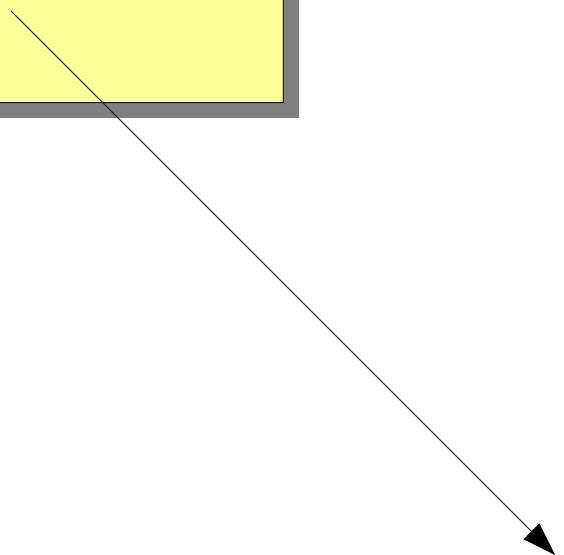
Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);
int mymax = mymin + n/nprocs -1;
while (!done) do
    mydiff = diff = 0;

    for i ← mymin to mymax
        for j ← 1 to n do
            /* avg neighbours */
            /* accumulate mydiff */
        endfor
    endfor
    LOCK(diff_lock);
    diff += mydiff;
    UNLOCK(diff_lock);
    BARRIER (bar1, nprocs);
    if (diff < TOL) then done = 1;
    BARRIER (bar1, nprocs);
endwhile
```

What happens if one thread runs ahead and executes



Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);
int mymax = mymin + n/nprocs -1;
while (!done) do
    mydiff = diff = 0;

    for i ← mymin to mymax
        for j ← 1 to n do
            /* avg neighbours */
            /* accumulate mydiff */
        endfor
    endfor
    LOCK(diff_lock);
    diff += mydiff;
    UNLOCK(diff_lock);
    BARRIER (bar1, nprocs);
    if (diff < TOL) then done = 1;
    BARRIER (bar1, nprocs);
endwhile
```

What happens if one thread runs ahead and executes while other threads are executing

Shared Address – Ocean Kernel

```
procedure Solve(A)
```

```
int mymin = 1 + (pid * n/procs);
int mymax = mymin + n/nprocs -1;
while (!done) do
    mydiff = diff = 0;

    for i ← mymin to mymax
        for j ← 1 to n do
            /* avg neighbours */
            /* accumulate mydiff */
        endfor
    endfor
    LOCK(diff_lock);
    diff += mydiff;
    UNLOCK(diff_lock);
    BARRIER (bar1, nprocs);
    if (diff < TOL) then done = 1;
    BARRIER (bar1, nprocs);
endwhile
```

What happens if one thread runs ahead and executes while other threads are executing

All threads have begin for loops after initializing diff, mydiff

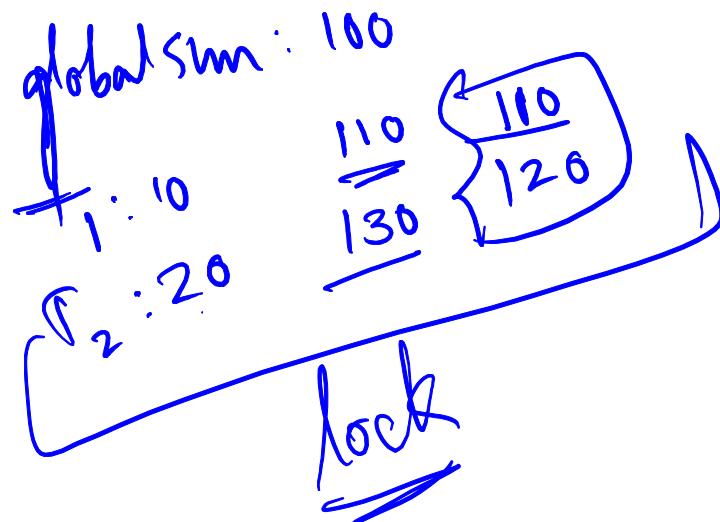
Shared Address – Ocean Kernel

procedure Solve(A)

```
int mymin = 1 + (pid * n/procs);
int mymax = mymin + n/nprocs -1;
while (!done) do
    mydiff = diff = 0; ←
    BARRIER(bar1,nprocs); ←
    for i ← mymin to mymax
        for j ← 1 to n do
            /* avg neighbours */
            /* accumulate mydiff */
        endfor
    endfor
    LOCK(diff_lock);
    diff += mydiff;
    UNLOCK(diff_lock);
    BARRIER (bar1, nprocs);
    if (diff < TOL) then done = 1;
    BARRIER (bar1, nprocs); ←
endwhile
```

What happens if one thread runs ahead and executes while other threads are executing

All threads have begin for loops after initializing diff, mydiff



Shared Address – Ocean Kernel

```
procedure Solve(A)
    int i, j, pid, done=0;
    float temp, mydiff=0;
    int mymin = 1 + (pid * n/procs);
    int mymax = mymin + n/nprocs - 1;
    while (!done) do
        mydiff = diff = 0;
        BARRIER(bar1,nprocs);
        for i ← mymin to mymax
            for j ← 1 to n do
                ...
            endfor
        endfor
        LOCK(diff_lock);
        diff += mydiff;
        UNLOCK(diff_lock);
        BARRIER (bar1, nprocs);
        if (diff < TOL) then done = 1;
        BARRIER (bar1, nprocs);
    endwhile
```

Shared Address – Ocean Kernel

```
main()
begin
    read(n); read(nprocs);
    A ← G_MALLOC();
    initialize (A);
    CREATE (nprocs,Solve,A);
    WAIT_FOR_END (nprocs);
end main
```

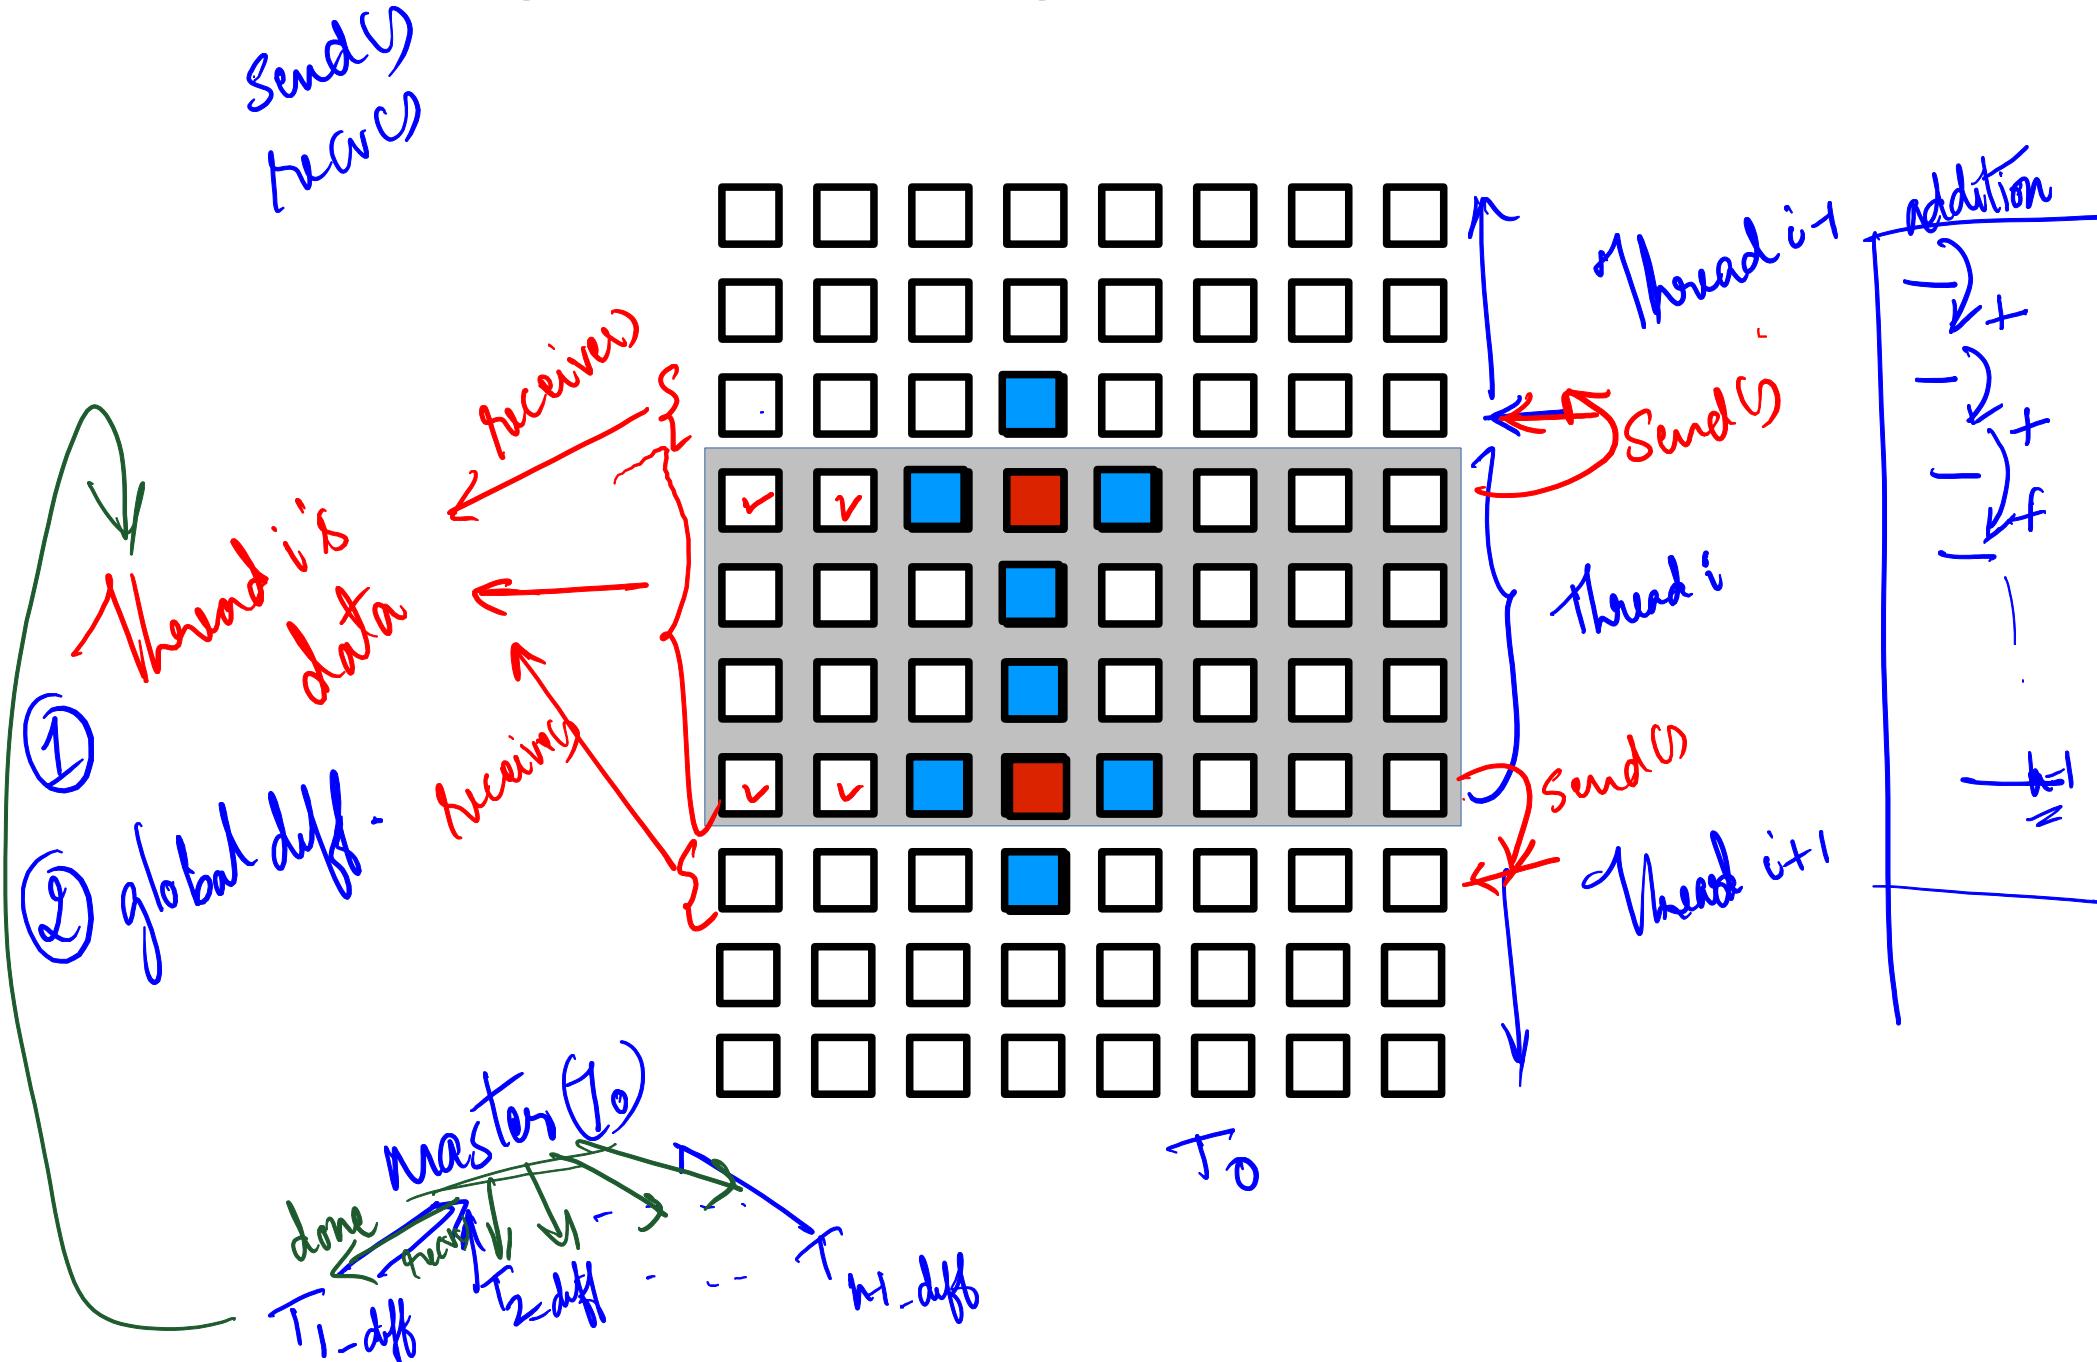
```
procedure Solve(A)
    int i, j, pid, done=0;
    float temp, mydiff=0;
    int mymin = 1 + (pid * n/procs);
    int mymax = mymin + n/nprocs - 1;
    while (!done) do
        mydiff = diff = 0;
        BARRIER(bar1,nprocs);
        for i ← mymin to mymax
            for j ← 1 to n do
                ...
            endfor
        endfor
        LOCK(diff_lock);
        diff += mydiff;
        UNLOCK(diff_lock);
        BARRIER (bar1, nprocs);
        if (diff < TOL) then done = 1;
        BARRIER (bar1, nprocs);
    endwhile
```

Shared Address – Ocean Kernel

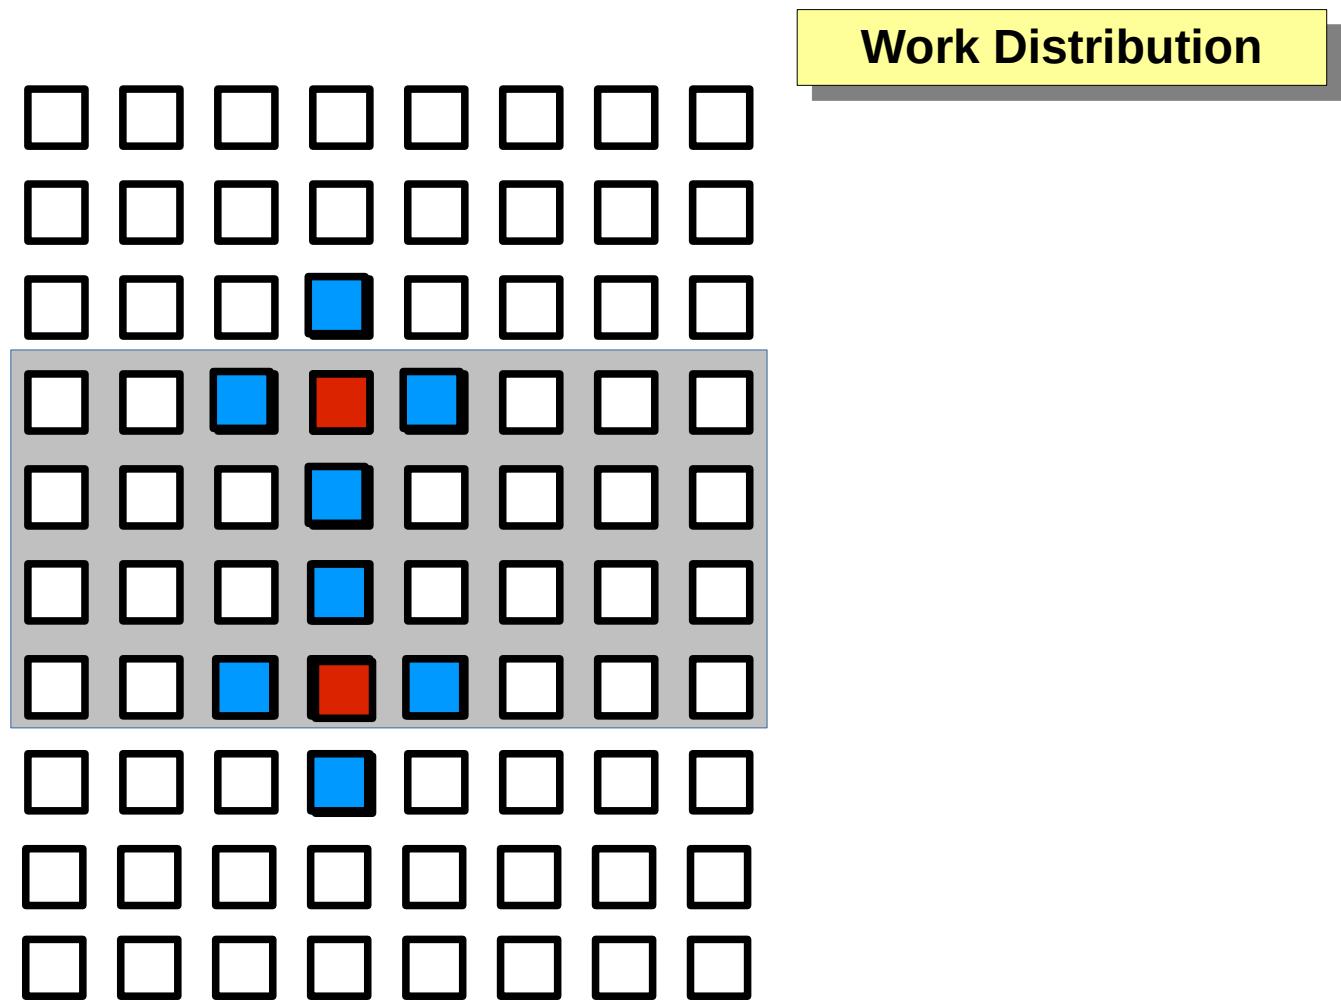
```
int n, nprocs;  
float **A, diff;  
LOCKDEC(diff_lock);  
BARDEC(bar1);  
  
main()  
begin  
    read(n); read(nprocs);  
    A ← G_MALLOC();  
    initialize (A);  
    CREATE (nprocs,Solve,A);  
    WAIT_FOR_END (nprocs);  
end main
```

```
procedure Solve(A)  
    int i, j, pid, done=0;  
    float temp, mydiff=0;  
    int mymin = 1 + (pid * n/procs);  
    int mymax = mymin + n/nprocs - 1;  
    while (!done) do  
        mydiff = diff = 0;  
        BARRIER(bar1,nprocs);  
        for i ← mymin to mymax  
            for j ← 1 to n do  
                ...  
            endfor  
        endfor  
        LOCK(diff_lock);  
        diff += mydiff;  
        UNLOCK(diff_lock);  
        BARRIER (bar1, nprocs);  
        if (diff < TOL) then done = 1;  
        BARRIER (bar1, nprocs);  
    endwhile
```

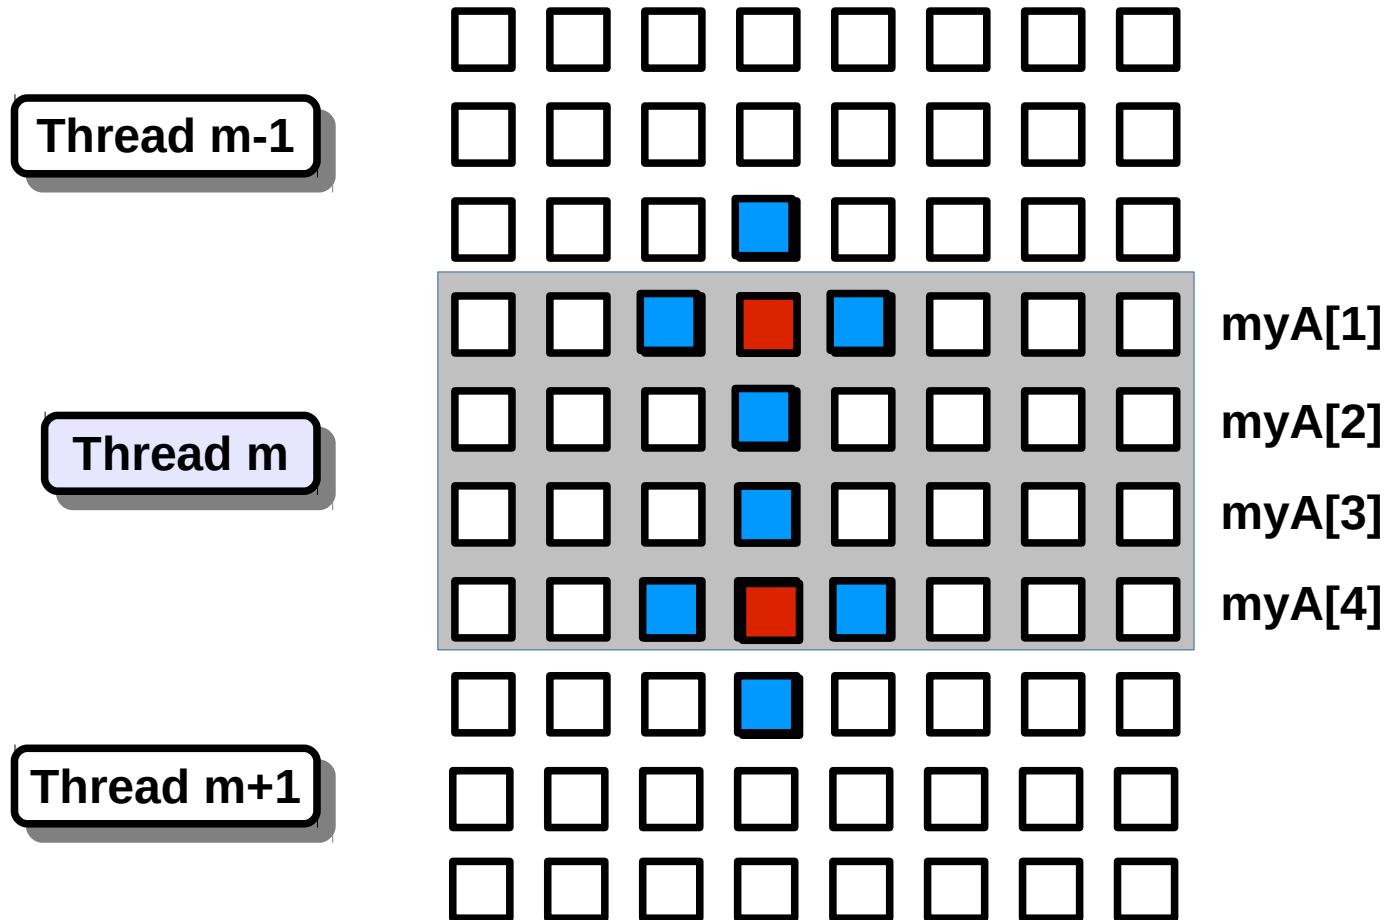
Message Passing – Ocean Kernel



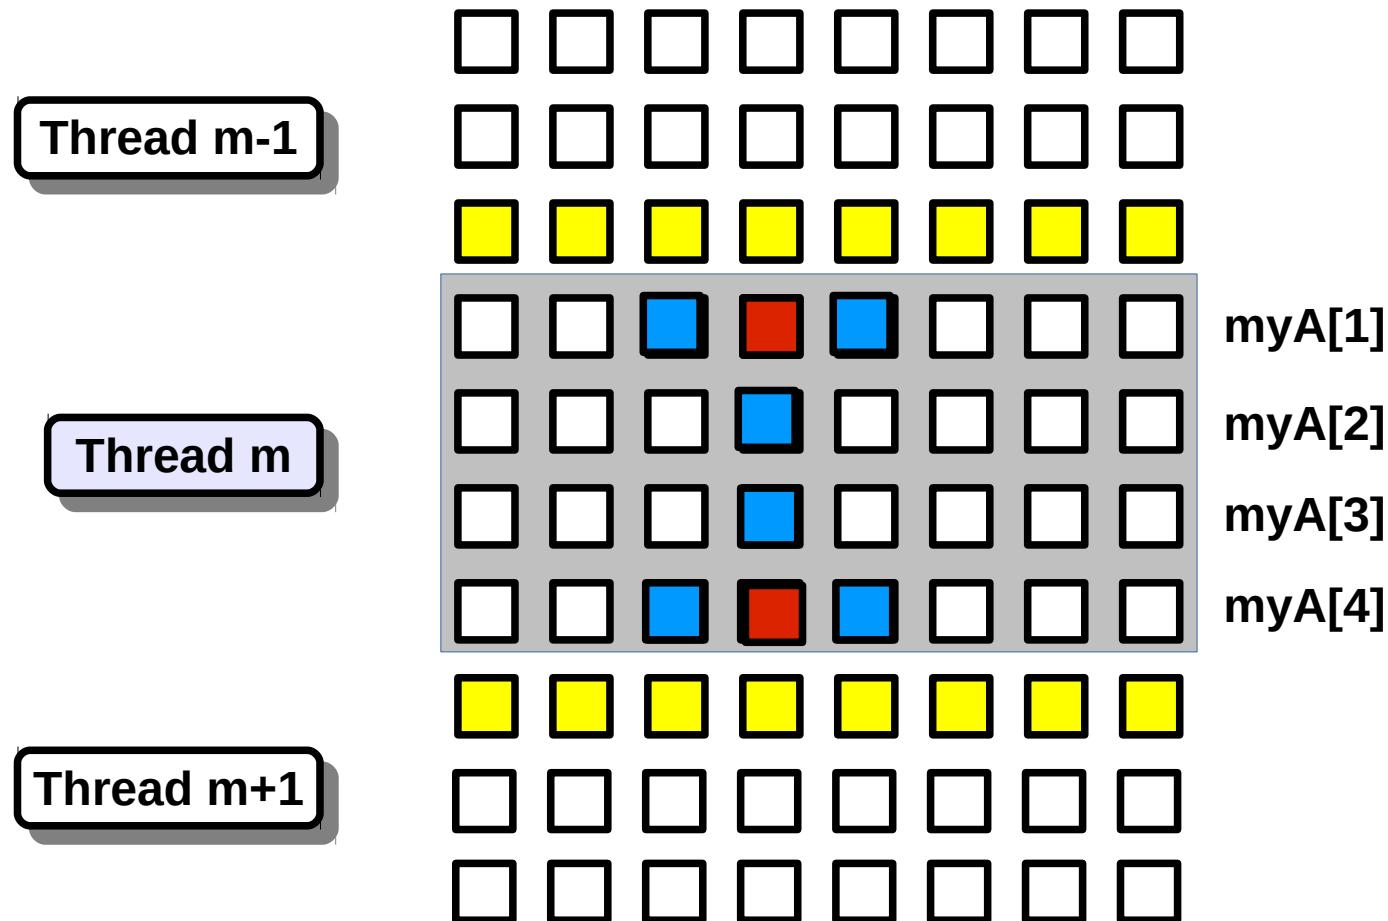
Message Passing – Ocean Kernel



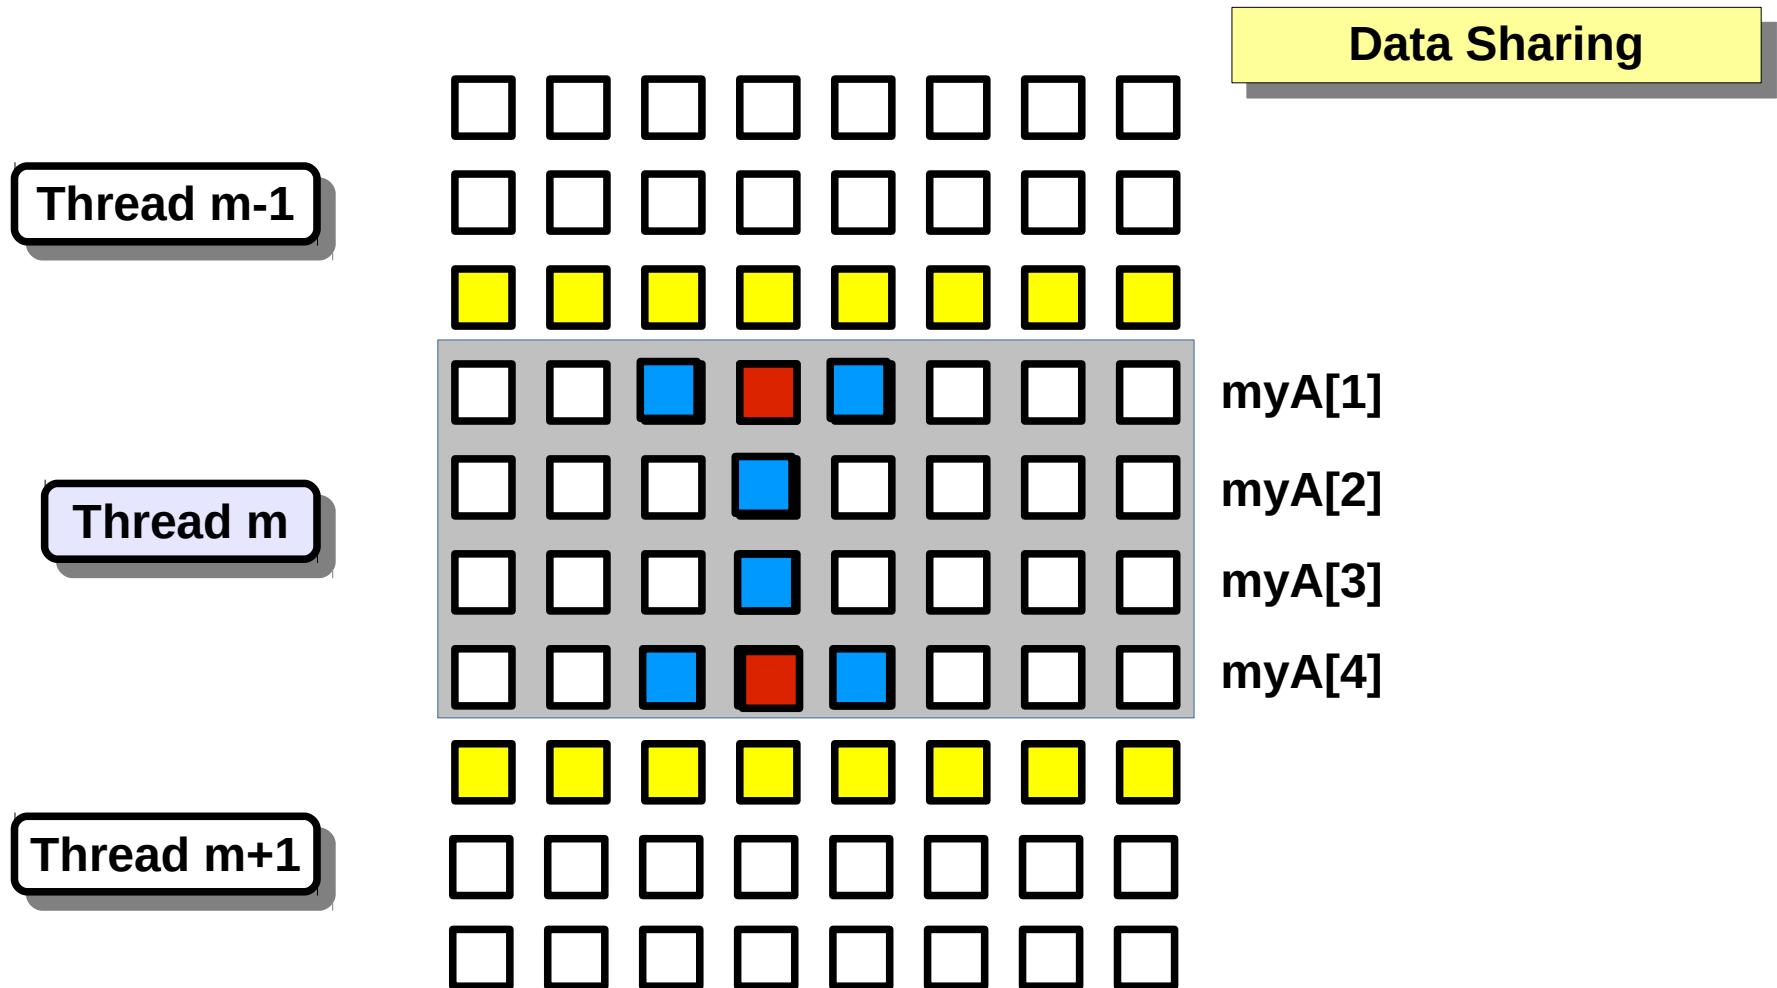
Message Passing – Ocean Kernel



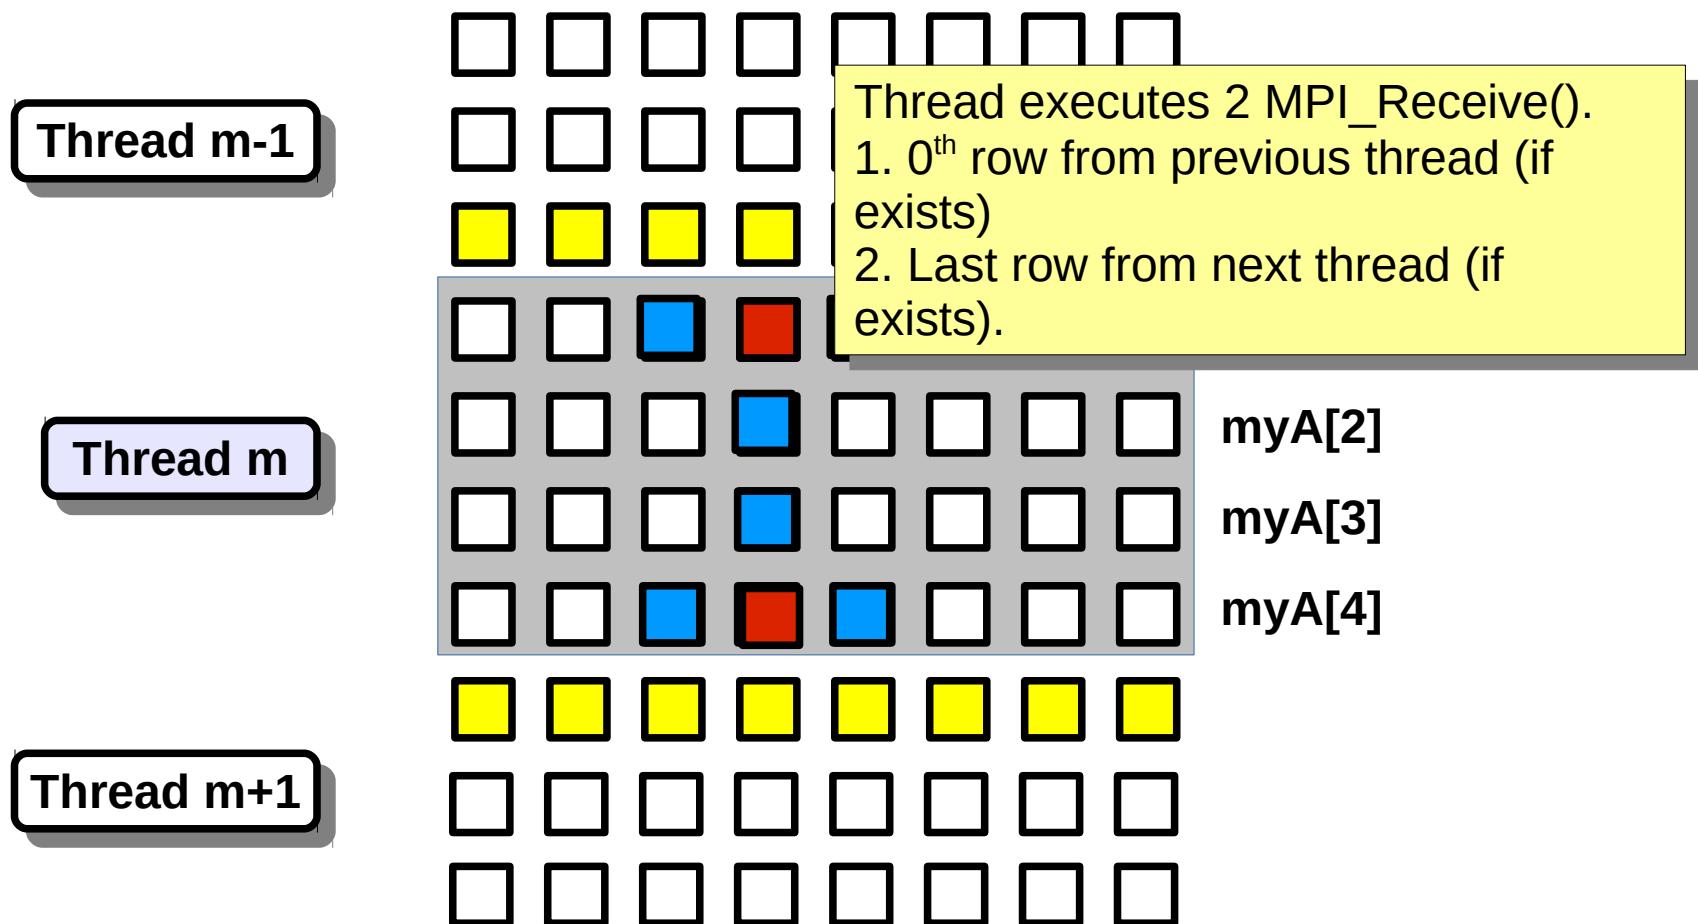
Message Passing – Ocean Kernel



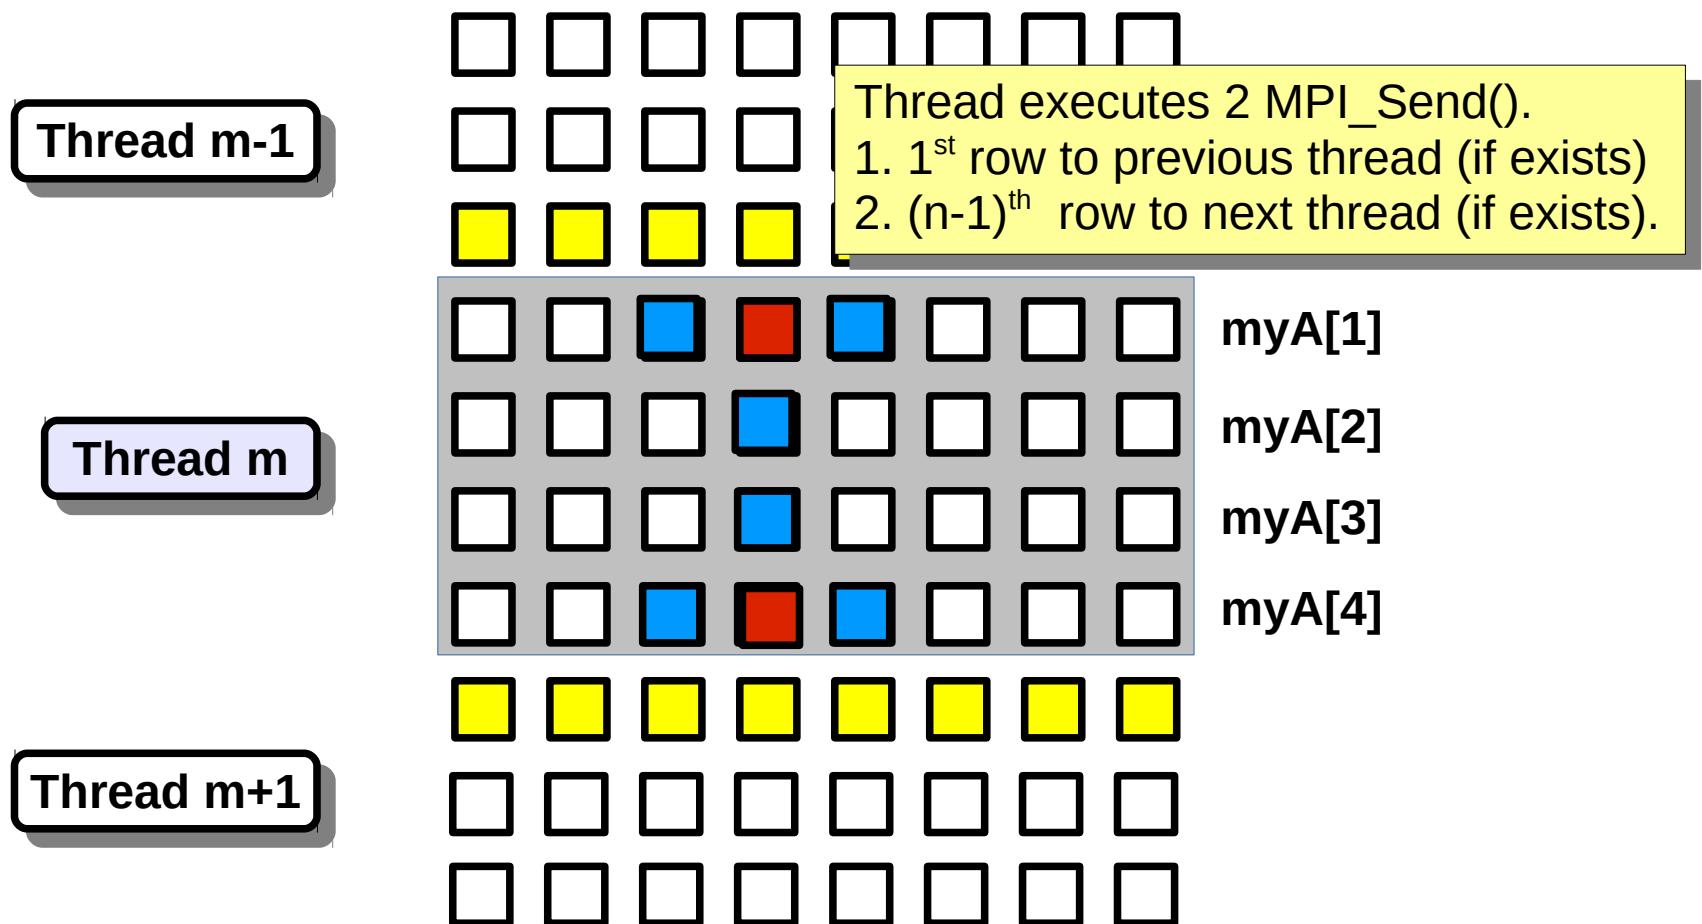
Message Passing – Ocean Kernel



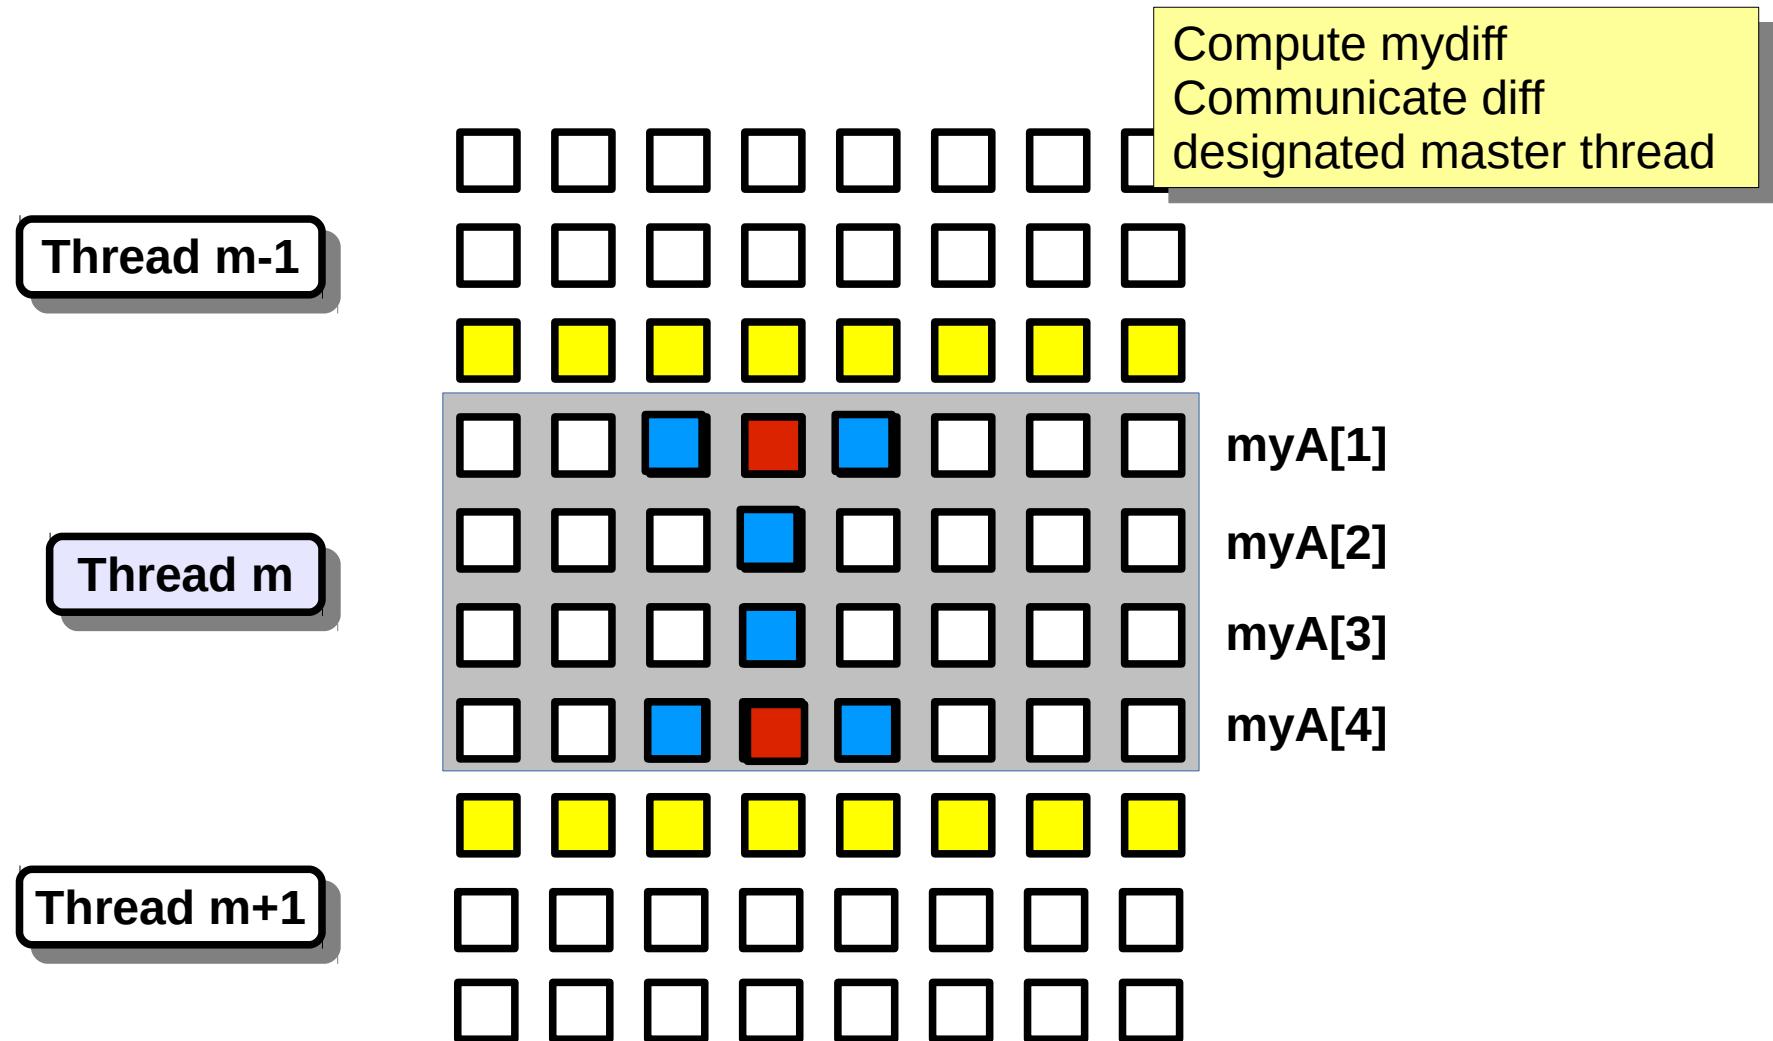
Message Passing – Ocean Kernel



Message Passing – Ocean Kernel



Message Passing – Ocean Kernel



Message Passing – Ocean Kernel

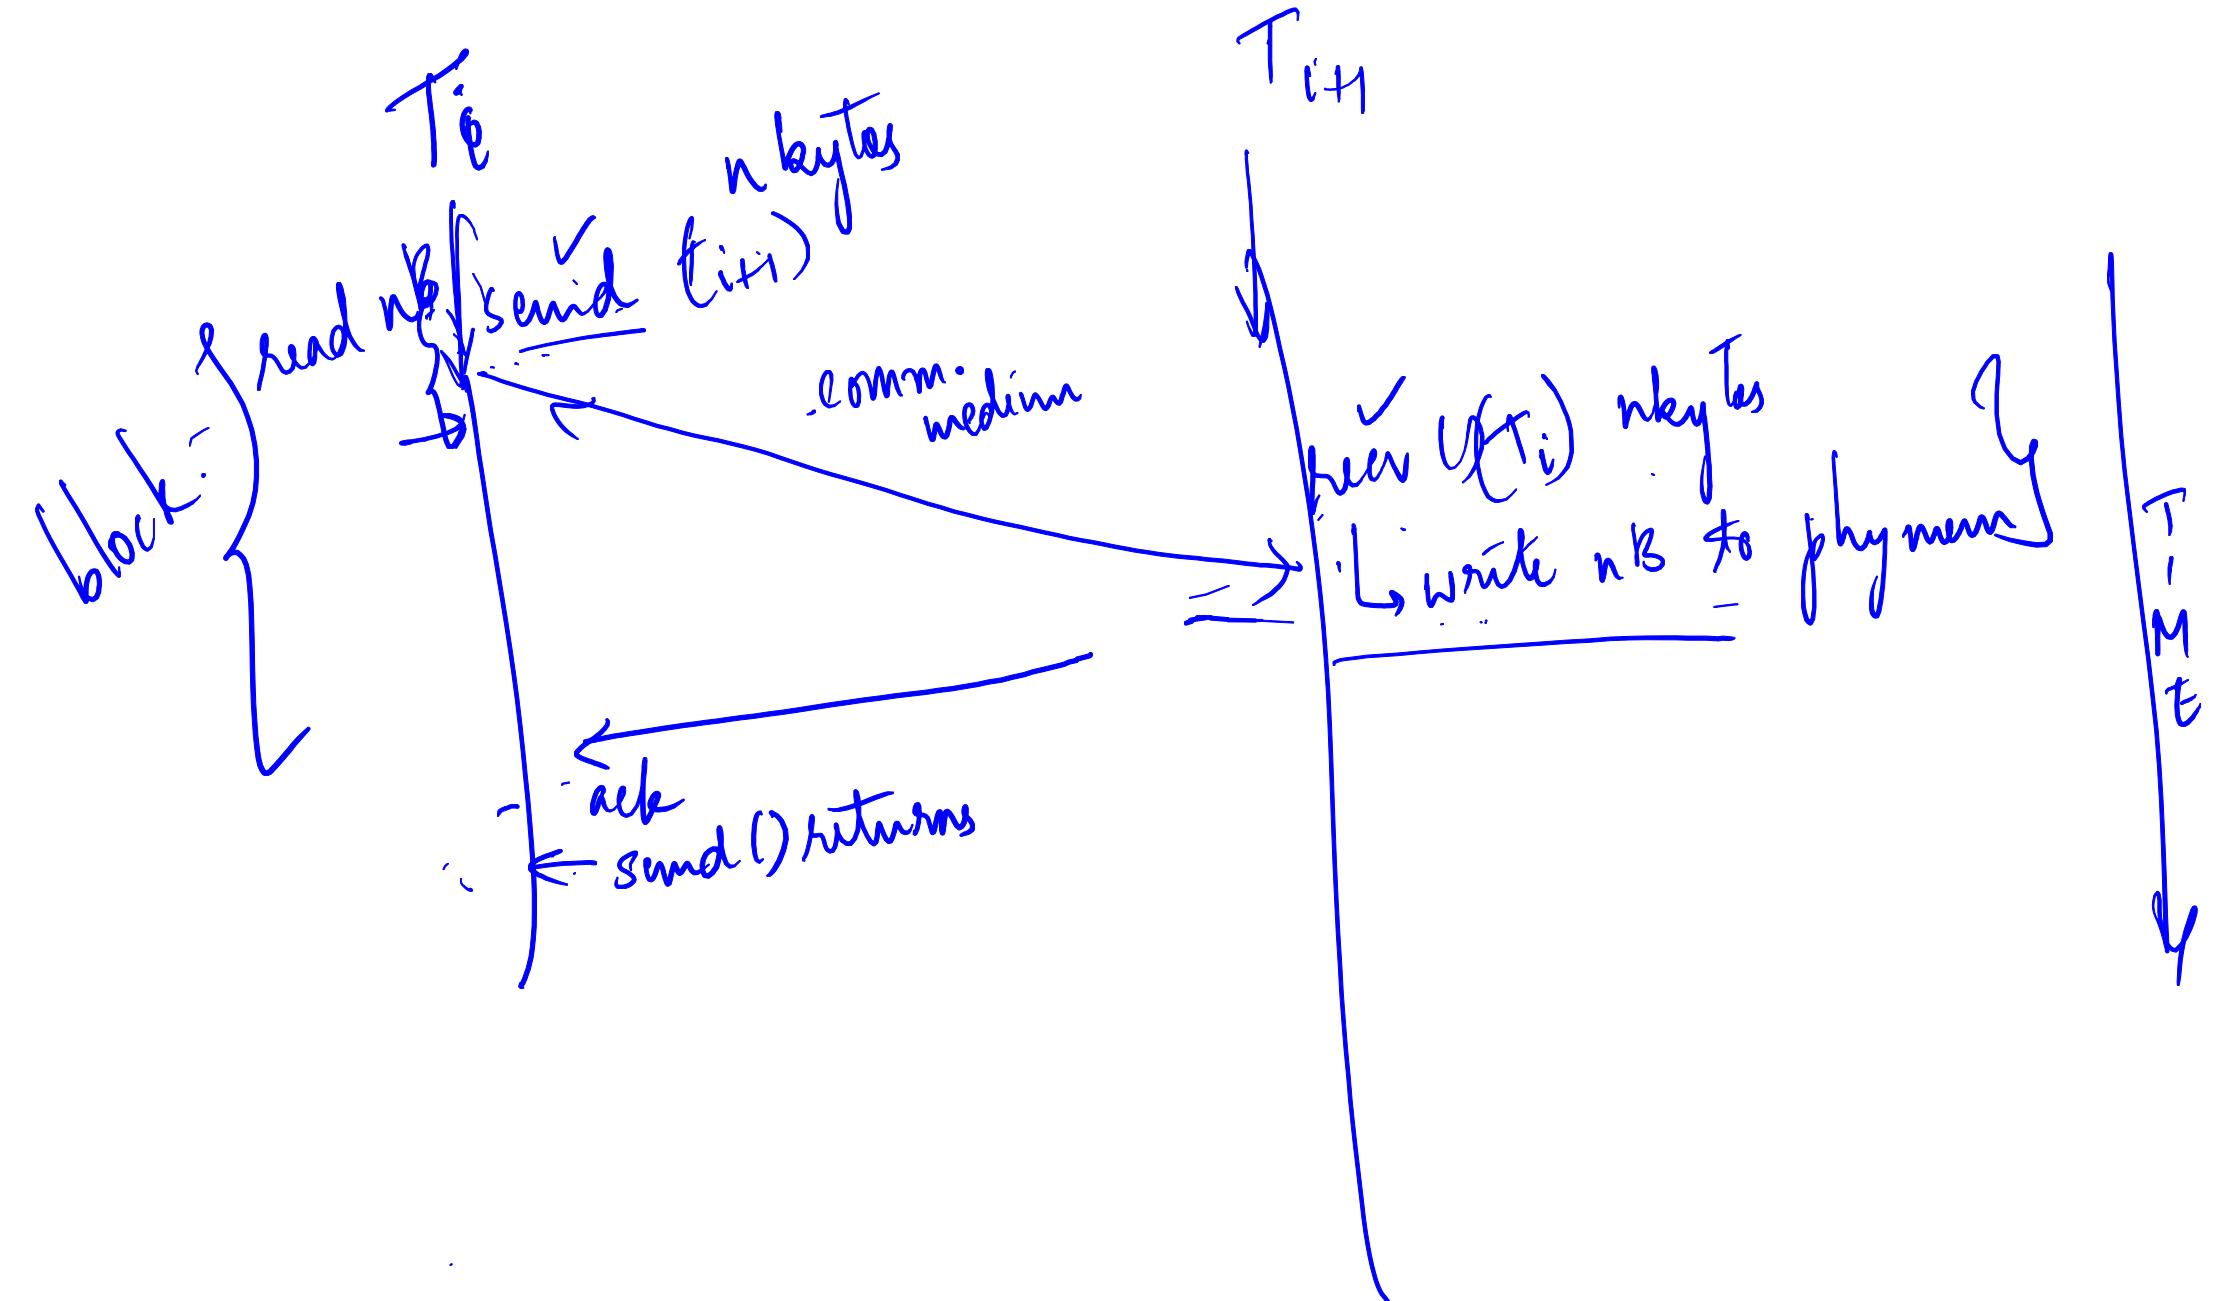
procedure Solve()

input data

if (pid != 0) *first row* *Thread id < pid*
SEND(&myA[1,0], n, pid-1, ROW);
if (pid != nprocs-1)
SEND(&myA[nn,0], n, pid+1, ROW);
if (pid != 0)
RECEIVE(&myA[0,0], n, pid-1, ROW);
if (pid != nprocs-1)
RECEIVE(&myA[nn+1,0], n, pid+1, ROW);

read
write
odd row
elements
communicating Thread id
size *how many elements?*
 sizeof(element);

** blocking*
** non-blocking*



Message Passing – Ocean Kernel

```
procedure Solve()
```

Compute mydiff

```
if (pid != 0)
    SEND(&myA[1,0], n, pid-1, ROW);
if (pid != nprocs-1)
    SEND(&myA[nn,0], n, pid+1, ROW);
if (pid != 0)
    RECEIVE(&myA[0,0], n, pid-1, ROW);
if (pid != nprocs-1)
    RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```

Message Passing – Ocean Kernel

```
procedure Solve()  
  
if (pid != 0)  
    SEND(&myA[1,0], n, pid-1, ROW);  
if (pid != nprocs-1)  
    SEND(&myA[nn,0], n, pid+1, ROW);  
if (pid != 0)  
    RECEIVE(&myA[0,0], n, pid-1, ROW);  
if (pid != nprocs-1)  
    RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```



```
for i ← 1 to nn do  
    for j ← 1 to n do  
        ...  
    endfor  
endfor
```

Message Passing – Ocean Kernel

```
procedure Solve()  
  
if (pid != 0)  
    SEND(&myA[1,0], n, pid-1, ROW);  
if (pid != nprocs-1)  
    SEND(&myA[nn,0], n, pid+1, ROW);  
if (pid != 0)  
    RECEIVE(&myA[0,0], n, pid-1, ROW);  
if (pid != nprocs-1)  
    RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```

```
for i ← 1 to nn do  
    for j ← 1 to n do  
        ...  
    endfor  
endfor
```

Communicate mydiff to a
designated master

Message Passing – Ocean Kernel

```
procedure Solve()  
  
if (pid != 0)  
    SEND(&myA[1,0], n, pid-1, ROW);  
if (pid != nprocs-1)  
    SEND(&myA[nn,0], n, pid+1, ROW);  
if (pid != 0)  
    RECEIVE(&myA[0,0], n, pid-1, ROW);  
if (pid != nprocs-1)  
    RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```

```
for i ← 1 to nn do  
    for j ← 1 to n do  
        ...  
    endfor  
endfor  
  
SEND(mydiff, 1, 0, DIFF);
```



Message Passing – Ocean Kernel

```
procedure Solve()  
  
if (pid != 0)  
    SEND(&myA[1,0], n, pid-1, ROW);  
if (pid != nprocs-1)  
    SEND(&myA[nn,0], n, pid+1, ROW);  
if (pid != 0)  
    RECEIVE(&myA[0,0], n, pid-1, ROW);  
if (pid != nprocs-1)  
    RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```

```
for i ← 1 to nn do  
    for j ← 1 to n do  
        ...  
    endfor  
endfor  
  
SEND(mydiff, 1, 0, DIFF);
```

Master Thread:
1. accumulate all mydiff
2. Check threshold

Message Passing – Ocean Kernel

```
procedure Solve()
```

```
if (pid != 0)
    SEND(&myA[1,0], n, pid-1, ROW);
if (pid != nprocs-1)
    SEND(&myA[nn,0], n, pid+1, ROW);
if (pid != 0)
    RECEIVE(&myA[0,0], n, pid-1, ROW);
if (pid != nprocs-1)
    RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```

```
for i ← 1 to nn do
    for j ← 1 to n do
        ...
    endfor
endfor
if (pid != 0)
    SEND(mydiff, 1, 0, DIFF);
else
    for i ← 1 to nprocs-1 do
        RECEIVE(tempdiff, 1, *, DIFF);
        mydiff += tempdiff;
    endfor
    if (mydiff < TOL) done = 1;
endwhile
```

all threads other than 0

mydiff

↑ {

Message Passing – Ocean Kernel

```
procedure Solve()
```

```
if (pid != 0)
    SEND(&myA[1,0], n, pid-1, ROW);
if (pid != nprocs-1)
    SEND(&myA[nn,0], n, pid+1, ROW);
if (pid != 0)
    RECEIVE(&myA[0,0], n, pid-1, ROW);
if (pid != nprocs-1)
    RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```

```
for i ← 1 to nn do
    for j ← 1 to n do
        ...
    endfor
endfor
if (pid != 0)
    SEND(mydiff, 1, 0, DIFF);
```

Receive result from master

```
else
    for i ← 1 to nprocs-1 do
        RECEIVE(tempdiff, 1, *, DIFF);
        mydiff += tempdiff;
    endfor
    if (mydiff < TOL) done = 1;
```

Communicate result to all threads

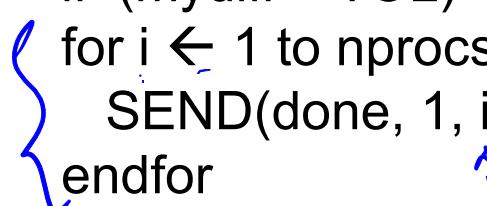
```
endwhile
```

Message Passing – Ocean Kernel

procedure Solve()

```
if (pid != 0)
    SEND(&myA[1,0], n, pid-1, ROW);
if (pid != nprocs-1)
    SEND(&myA[nn,0], n, pid+1, ROW);
if (pid != 0)
    RECEIVE(&myA[0,0], n, pid-1, ROW);
if (pid != nprocs-1)
    RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```

```
for i ← 1 to nn do
    for j ← 1 to n do
        ...
    endfor
endfor
if (pid != 0)
    SEND(mydiff, 1, 0, DIFF);
→ RECEIVE(done, 1, 0, DONE);
else
    for i ← 1 to nprocs-1 do
        RECEIVE(tempdiff, 1, *, DIFF);
        mydiff += tempdiff;
    endfor
    if (mydiff < TOL) done = 1;
    for i ← 1 to nprocs-1 do
        SEND(done, 1, i, DONE);
    endfor
endif
endwhile
```



Message Passing – Ocean Kernel

```
procedure Solve()
    int i, j, pid, nn = n/nprocs, done=0;
    float temp, tempdiff, mydiff = 0;
    myA ← malloc(...)

    initialize(myA);

    while (!done) do
        mydiff = 0;
        if (pid != 0)
            SEND(&myA[1,0], n, pid-1, ROW);
        if (pid != nprocs-1)
            SEND(&myA[nn,0], n, pid+1, ROW);
        if (pid != 0)
            RECEIVE(&myA[0,0], n, pid-1, ROW);
        if (pid != nprocs-1)
            RECEIVE(&myA[nn+1,0], n, pid+1, ROW);
```

```
for i ← 1 to nn do
    for j ← 1 to n do
        ...
    endfor
endfor
if (pid != 0)
    SEND(mydiff, 1, 0, DIFF);
    → RECEIVE(done, 1, 0, DONE);
else
    for i ← 1 to nprocs-1 do
        → RECEIVE(tempdiff, 1, *, DIFF);
        mydiff += tempdiff;
    endfor
    if (mydiff < TOL) done = 1;
    for i ← 1 to nprocs-1 do
        SEND(done, 1, i, DONE);
    endfor
endif
endwhile
```

