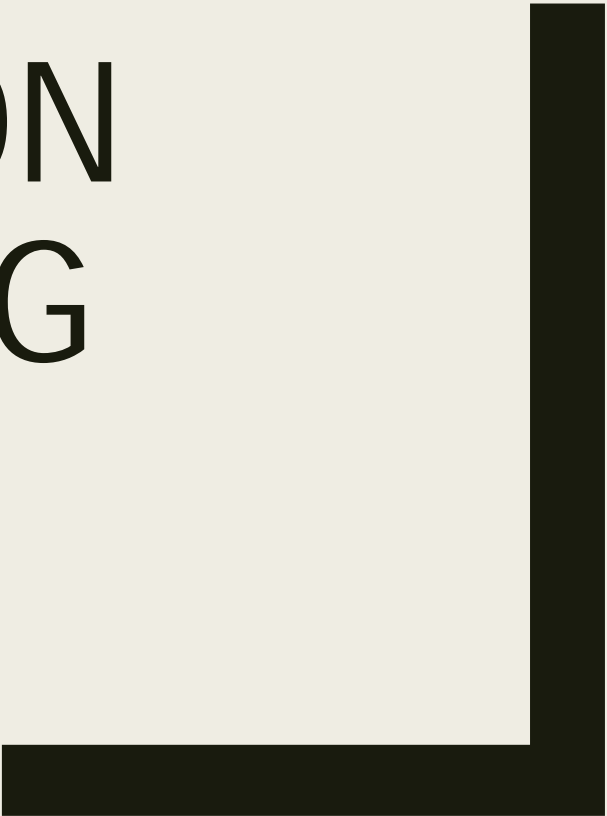





TRANSACTION PROCESSING

Dr.M.Venkatesan
Assistant Prof,CSE,NITK



Transaction Management

- A transaction is a program unit whose execution may change contents of a database.



- Transaction is used to represent a logical unit of database processing that must be completed in its entirety to ensure correctness.

Example:

Bank(db) – 10,000
Debit 5000

R(a) → current
A-5000
W(a) } Transaction

ACID properties of Transaction

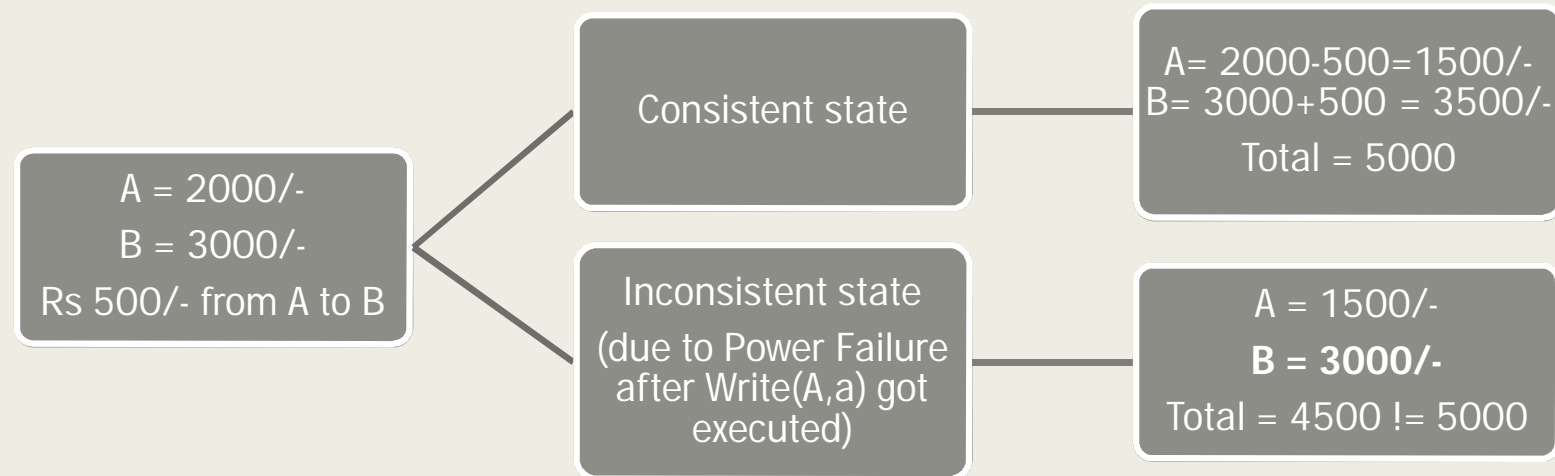
■ Atomicity(ALL or None):

Ensure that a transaction will run to completion as an individual unit, at the end of which either no changes have occurred to the database, or database has been changed in a consistent manner.

Example:

Transaction(Ti):

Read(A,a)
A – 500
Write(A,a)
Read(B,b)
B = b + 500
Write(B,b)



Cont'd..

- **Consistency (Correctness):**

Ensures that if the database was in a consistent state before the start of the transaction, then or termination, the database will also be in a consistent state.

Sum(A,B) before
transaction



Sum(A,B) after
transaction

Cont'd..

- **Isolation:**

Indicates that the actions performed by a transaction will be isolated or hidden from outside the transaction until it terminates.

Transaction T1	Transaction T2
R(A) A – 500 W(A)	
	R(A) R(B) Print(a,b) → 4,500
R(B) B = B + 500 W(B)	

Initially
A = 1500
B = 3000

- **Durability(Commit):**

All updates done by a Transaction must become permanent (or)

Ensures that the commit action of a transaction on its termination will be reflected in the table

Concurrent Execution:

- It implies interleaving execution of operations of a transaction.
- Benefits
 - i. Helps in reducing waiting time.
 - ii. Improved throughput and resource utilization.

Schedule:

It represents the order in which instructions of a transaction are executed

T1 → T2 → T3  schedule

Problems with Concurrent Execution

- Lost update Problem(W – W conflict)

Occurs when two transactions that accesses the same database items have their operations interleaved in a way that makes the value of the database item incorrect.

Update is lost	{	Transaction (T1)	Transaction(T2)
		R(A) → 1000 A = A -50 → 950	
			R(A) → 1000 A = A + 100 → 1100
		W(A) → 950	
			W(A) → 1100

Actual value: $1000 - 50 = 950$
 $950 + 100 = 1050$

Cont'd:

■ Temporary update(Dirty Read) Problem [W – R Conflict]

Occurs when one transaction updates a database item and then the transaction fails, but its update is read by some other transaction.

A = 100

Power
failure
Roll Back

Transaction T1	Transaction T2
R(A) → 100 A = A + 20 → 120 W(A)	
	R(A) → 120 A = A + 10 → 130 W(A) → 130
R(B)	
	Commit (130) (110)

Updated value

Cont'd:

■ Unrepeatable Read[W – R conflict]

If a transaction 'Ti' reads an item value twice and the item is changed by another transaction 'Tj' in between the two read operation. Hence 'Ti' receives different values for its two Read operation of the same item.

A = 1000

Transaction T1	Transaction T2
R(A) → 1000	
	R(A) A = A + 2000, A = A – 2000 W(A)
R(A) → 3000	

Cont'd:

■ Incorrect Summary Problem:

If one transaction is calculating an aggregate summary function on a no. of records, the aggregate function may calculate some values before they are updated and others after they are updated – results in incorrect summary.

A = 100, Y = 200

Transaction T1	Transaction T2
	Sum = 0 R(A) Sum = Sum + A , Sum = 100 R(Y) Sum = Sum + Y, Sum = 300
R(Y) → 200 Y = Y + 100 → 300	



Schedule

- A schedule 's' of n transactions T_1, T_2, \dots, T_n is an ordering of operations of the transactions in chronological order.

$T_i(x \rightarrow y \text{ operation})T_j$

In any schedule $T_i(x) \rightarrow T_j(y)$

T1	T2
R(X)	
	W(X)

$T_1 \rightarrow T_2$

When several transactions are executing concurrently then the order of execution of various instructions is known as Schedule.

Types of Schedule

■ Serial Schedule:

Does not interleave the actions of any operations of different transactions.

Always ensure a Consistent state.

INITIALLY $A = 100$

T1	T2
R(A) $A = A + 50$ W(A)	
	R(A) $A = A + 100$ W(A)

T1 \rightarrow T2

T1: $A = 100 + 50 = 150$

T2: $A = 150 + 100 = 250$

T2 \rightarrow T1

T2: $A = 100 + 100 = 200$

T1: $A = 200 + 50 = 250$

Non - Serial

T1	T2
R(A) $A = A + 50$	
⋮	R(A) $A = A + 100$
W(A) $\rightarrow 150$	⋮
	W(A) $\rightarrow 200$

Cont'd:

■ Complete Schedule:


If the last operation of each transaction is either abort or commit.

T1	T2	COMPLETE SCHEDULE
R(A)		
	R(A)	
W(A)		
COMMIT		
	W(A) ABORT	

Cont'd:

■ Recoverable Schedule:

Is one where for each pair of transactions (T_i , T_j), such that T_j reads a data item that was previously written by T_i , then the commit operation of T_i should appear before commit operation of T_j .



Cont'd..

■ Cascade less Schedule:

Is one where for each pair of transaction(T_i, T_j) such that T_j reads a data item written by T_i , then the commit operation of T_i should appear the read operation of T_j .

Cascading rollback

T1	T2	T3
R(A) W(A)		
	R(A) W(A)	
		R(A) W(A)
COMMIT		
	COMMIT	
		COMMIT

T1	T2	T3
R(A) W(A) COMMIT		
	R(A) W(A) COMMIT	
		R(A) W(A) COMMIT

fail

Cont'd:

■ Strict Schedule:

If a value written by a transaction cannot be read or overwritten by another transaction until the transaction is either aborted or committed.

T1	T2
R(A) W(A) COMMIT	
	W(A) R(A)

Every Strict Schedule is
both Recoverable and
cascade less.

Serializability

- A schedule 's' of 'n' transaction is serializable if it is equivalent to some serial schedule of the same 'n' transactions.

- **Conflict Serializable:**

If it is conflict equivalent to serial schedule

R1A → W2A
W1A → R2A
W2A

R1A → W2A

T1	T2
R(A) W(A)	
	R(A) W(A)
R(B) W(B)	
	R(B) W(B)



T1	T2
R(A) W(A) R(B) W(B)	
	R(A) W(A) R(B) W(B)

$S1 \stackrel{C}{=} S2$

Test for Conflict Serializability


- Precedence graph is used
- Let 's' be a schedule, construct a directed graph known as precedence graph.
- Graph consists of a pair of $G = (V, E)$ where V : a set of vertices, E : a set of edges
- Algorithm for creating graph:
 - Create a node for each transaction
 - A directed edge, $T_i \rightarrow T_j$, if T_j reads a value of an item written by T_i .
 - Directed edge $T_i \rightarrow T_j$, if T_j writes a value into item after it has been read by T_i .
 - Directed edge, $T_i \rightarrow T_j$, if T_j write after T_i write.

A schedule is conflict serializable if and only if precedence graph is acyclic




Check for Conflict Serializable

T1	T2	T3
R(X)		
	R(Z)	
R(Z)		
		R(X) R(Y) W(X)
	R(Y) W(Z) W(Y)	



Check for conflict Serializability

T1	T2	T3
R(X)		
		R(X) W(X)
W(X)		
	R(X)	



Non conflict serializable

View Serializability

- Two schedules 's' and 's1' are view equivalent if the following conditions are met:
 - i. For each data item Q, if T_i reads an initial value of Q in schedule S, then T_i must in s1 also reads an initial value of Q
 - ii. If T_i executes reads Q in s, and that value was produced by T_j (if any), then T_i must in schedule s1 also reads the value of Q that was produced by T_j .
 - iii. For each data item Q, the transaction that perform the final write(Q) operation in schedule s must also perform the final write(Q) in schedule s1.

A schedule is view serializable if it is view equivalent to a serial schedule

Example of View Serializability

s1		s2		s3	
T1	T2	T1	T2	T2	T1
R(A) W(A)		R(A) W(A) R(B) W(B)		R(A) W(A) R(B) W(B)	
	R(A) W(A)		R(A) W(A) R(B) W(B)		R(A) W(A) R(B) W(B)
R(B) W(B)					
	R(B) W(B)				

$$S1 \stackrel{V}{=} S2$$

$$S1 \not\stackrel{V}{=} S3$$

- **Conflict Operations**: operations are said to be conflicting if
 - i. Belong to different transactions
 - ii. Access to same database item 'A'.
 - iii. Atleast one of them is a write operation \rightarrow R W, W R, W W

R R – NOT
CONFLICT

T1	T2
R(A)	
	R(A)
W(A)	

S1:
 $R(A) \rightarrow 100$
 $A = A + 10 \rightarrow 110$
 $W(A)$

S2:
 $R(A) \rightarrow 100$
 $A * 1.1 \rightarrow 110$
 $W(A)$

- **Equivalent Schedule**: Two schedules 's1' and 's2' are said to be equivalent schedule if they produce the same final database state.

*) **Result equivalent schedule**: Produce same final DB state for same initial values of data.

■ Conflict Equivalent:

Two schedules are said to be conflict equivalent if all conflicting operations in both the schedule must be executed in the same order.

Question1) check for conflict equivalent

S1: R1(A), R2(B), W1(A), W2(B)

S2: R2(B), R1(A), W2(B), W1(A)

S1		S2	
T1	T2	T1	T2
R(A)			R(B)
	R(B)	R(A)	
W(A)			W(B)
	W(B)	W(A)	

$S1 \stackrel{C}{=} S2$

Question2) check for conflict equivalent

S1: R1(A), W1(A), R2(B), W2(B), R1(B)

S2: R1(A), W1(A), R1(B), R2(B), W2(B)

S1 (W2B → R1B)		S2 (R1B → W2B)	
T1	T2	T1	T2
R(A) W(A)		R(A) W(A) R(B)	
	R(B) W(B)		R(B) W(B)
R(B)			

$S1 \not\stackrel{C}{=} S2$

